

# Revisiting the Compact Tail Assignment Model: Corrections and an Exact Full Maintenance Hierarchy Extension

Viktor Borodin<sup>a</sup>, Arun Kumar Pappala<sup>b</sup>

<sup>a</sup>*Taras Shevchenko National University, Kyiv, Ukraine*

<sup>b</sup>*Singapore University of Technology and Design, Singapore*

---

## Abstract

This paper revisits the compact mixed-integer linear programming model for the Tail Assignment Problem (TAP) introduced by Khaled et al. (2018). We make four contributions. First, we show that the basic continuity/turn-time criteria (C2–C4; model constraints (3)–(6) in Khaled et al.) admit a simultaneous-departure corner case that is mathematically feasible but physically infeasible, and we close this gap with explicit non-overlap constraints. Second, we identify and formally correct a flaw in the paper’s cumulative-flight-time constraint (Constraint 13): the original formulation allows maintenance intervals to be incorrectly relaxed. We show that a two-row endpoint split is also insufficient and give an exact cumulative-state formulation. Third, we extend the maintenance model beyond the type-A single-check framework to a full four-level  $\{A, B, C, D\}$  maintenance hierarchy with correct counter resets and pre-horizon accumulated-hours accounting. Fourth, we give a complete, day-free event-based reformulation of the same maintenance requirements, in which each flight’s own timestamp supplies all calendar information and the flight-hour counter is propagated exactly along

each aircraft’s actual route with no day-pair window, and we prove its correctness directly. We give a numeric counterexample, evaluated against the real built constraint objects at two adversarial check patterns, in which the original and endpoint-split formulations each fail on a different pattern while the event-based formulation rejects both, and we independently verify every reported solution by replaying its maintenance counters outside the model. We then report exact solves, cross-validated with two independent solvers, across targeted instances that isolate each check type (A, B, C, D) and their combination, demonstrating that the corrected hierarchy model enforces every check type without relaxation. All code, instances, and result tables accompany the project.

*Keywords:* Tail assignment, Aircraft routing, Maintenance scheduling, Integer programming, Mixed-integer programming corrections

---

## 1. Introduction

The airline planning process is traditionally decomposed into a sequence of sub-problems: timetable design, fleet assignment, crew rostering, tail assignment, and recovery planning (Barnhart et al., 2003; Gopalan and Talluri, 1998). The *tail assignment problem* (TAP)—assigning a specific aircraft (identified by its tail number) to each flight leg—is the penultimate planning step and is tightly coupled with aircraft maintenance scheduling. Poor tail assignment decisions directly inflate operational costs through unnecessary ferry (repositioning) flights, suboptimal maintenance routing, and safety-margin violations.

Khaled et al. (2018) proposed a compact 0–1 linear programming formulation for the TAP, featuring only  $O(n_f n_p)$  binary variables (where  $n_f$  is the number of flights and  $n_p$  the fleet size). This is a significant reduction over the classic multi-commodity flow model of Levin (1971), which requires  $O(n_f^2 n_p)$  variables. Computational experiments reported in that paper show that instances with up to 1 500 flight legs and 40 aircraft can be solved to certified optimality within one hour on a standard workstation.

Our non-overlap correction retains the  $O(n_f n_p)$  binary-variable space but can add  $O(n_f^2 n_p)$  conflict inequalities in the worst case. In practical timetables, only temporally conflicting pairs are instantiated; the resulting constraint count is therefore governed by the sparse conflict set  $\mathcal{O}$  defined in Section 3.

#### *Our contributions*

The present paper makes four contributions that check, correct, and extend those results:

1. **Basic feasibility correction (C2–C4 / constraints (3)–(6)).** We show that the strict-time balance definition used in the compact basic model admits a simultaneous-departure corner case: one aircraft can be assigned to two flights departing at the same airport and same time without violating the balance inequalities. We formalize this counterexample and enforce physical executability by adding explicit pairwise non-overlap constraints.
2. **Constraint (13) correction.** We show that the cumulative-flight-time maintenance constraint (Constraint 13 in Khaled et al. (2018))

contains a logical flaw in opening and one-endpoint intervals. We also show by counterexample that a natural two-row endpoint split does not fully repair the model. We therefore formulate an exact cumulative-state recursion that propagates and resets the flight-hour counter without day-pair ambiguity.

3. **Full  $\{A, B, C, D\}$  maintenance hierarchy.** The model of Khaled et al. (2018) handles type-A checks (cumulative flight hours). We extend the formulation to a four-level A/B/C/D hierarchy representing a traditional packaged maintenance program: in our model, type-A and B checks are flight-hour triggered and type-C and D checks are calendar-day triggered. A heavier check subsumes all lighter checks (hierarchical counter resets). We further account for pre-horizon accumulated maintenance hours/days per aircraft, and we validate the extension with a dedicated computational study that isolates each check type (A, B, C, D) and their combination, independently re-verifying every reported solution against the raw flight timetable.
4. **Event-based reformulation.** The day-indexed hierarchy model above still repeats a calendar-day index across every maintenance variable. We give a complete, day-free reformulation of the same C1–C14 requirements, in which each flight’s own timestamp supplies all calendar information and maintenance is represented as a virtual flight triggered by, and indexed on, the preceding real flight rather than by calendar day. This removes the day dimension from the maintenance model’s size and lets the flight-hour counter propagate exactly along each aircraft’s actual route, with no day-pair window and therefore no possibility of

the split-form failure mode proved for Constraint (13).

### *Paper organisation*

Section 2 reviews related work. Section 3 restates the basic TAP model and proves a C2–C4 corner-case correction. Section 4 presents the maintenance-extended model, the Constraint (13) correction, and the induced interpretation updates for C8–C14 under the corrected basic model. Section 5 describes the full  $\{A, B, C, D\}$  extension. Section 6 gives the day-free event-based reformulation of the same requirements and proves its correctness directly. Section 7 reports the numeric Constraint (13) counterexample and the per-check-type computational study. Section 8 concludes.

## **2. Literature Review**

### *2.1. Mathematical programming formulations*

Three main modelling paradigms underlie TAP research:

*Set partitioning / column generation..* The non-compact formulation enumerates *strings* (feasible flight sequences per aircraft) as 0–1 variables. Because the number of strings is exponential, column generation is required (Lavoie et al., 1988; Minoux, 1984; Barnhart et al., 1998; Sarac et al., 2006). Combining column generation with branch-and-bound (*branch-and-price*) gives exact solutions but at significant algorithmic overhead. Robust extensions appear in Borndörfer et al. (2010) and Froyland et al. (2013).

*Multi-commodity flow..* Levin (1971) introduced the compact multi-commodity flow formulation with  $O(n_f^2 n_p)$  binary variables. This remains the stan-

dard benchmark for compact models and is the primary comparison point in Khaled et al. (2018).

*Time-space networks.* Hane et al. (1995) proposed the time-space network for the *fleet assignment problem*. Nodes represent airport–time events; arcs encode ground waits, flights, and overnight stays. Extensions to aircraft routing appear in Clarke et al. (1996). Preprocessing by node aggregation and island isolation reduces model size (Hane et al., 1995; Sherali et al., 2006). Daily and weekly maintenance routing variants are studied in Liang et al. (2011) and Liang and Chaovalitwongse (2012); the latter also integrates fleet assignment decisions, but only approximate solutions are obtained via variable fixing.

*Nonlinear / lifted formulations.* Haouari et al. (2012) present a compact formulation with nonlinear constraints that are linearised via the Reformulation–Linearisation Technique (RLT). The largest instances solved involve 344 flights, an order of magnitude below the instances handled by the model of Khaled et al. (2018).

## 2.2. MILP alternatives to the Khaled compact model

Beyond the compact binary assignment model of Khaled et al. (2018), several mixed-integer alternatives are relevant for exact TAP modelling and large-scale implementation:

- **Arc-flow MILP (time-indexed or event-indexed):** decision variables are attached to feasible flight-connection arcs instead of flight–aircraft assignment pairs, often yielding stronger network-structure constraints (Levin, 1971; Hane et al., 1995; Sherali et al., 2006).

- **Route-based master MILP (set partitioning):** aircraft duties are represented as columns in a master MILP, typically solved with branch-and-price; this is exact but algorithmically heavier (Barnhart et al., 1998; Sarac et al., 2006; Klabjan, 2005).
- **Decomposed MILP frameworks:** Benders- or Dantzig–Wolfe-style decompositions separate assignment/routing from maintenance-resource feasibility, improving scalability on larger fleets (Cordeau et al., 2001; Mercier et al., 2005; Desaulniers et al., 1997).
- **Two-stage or rolling-horizon MILP:** first assign flights, then repair/optimize maintenance and turnaround feasibility in a second MILP, trading global optimality for tractable runtime (Clarke et al., 1996; Liang et al., 2011; Liang and Chaovalitwongse, 2012).
- **Robust/scenario MILP variants:** delay and disruption uncertainty is modelled through multiple scenarios with non-anticipativity-style linking constraints, producing schedules that are less brittle operationally (Borndörfer et al., 2010; Froyland et al., 2013).

In this taxonomy, the Khaled model is best viewed as a compact, directly solvable baseline that is simple to implement and benchmark, while the alternatives above trade model size, decomposition complexity, and robustness for stronger scalability or richer operational realism.

### *2.3. Comparison scope and fairness*

For an empirical comparison to be methodologically fair, competing models should be evaluated under a common instance family, objective definition,

maintenance semantics, and solver-budget protocol. In TAP, this condition is often not met across compact MILP, branch-and-price, Benders-decomposed, and robust/scenario formulations because these approaches typically optimize different decision scopes or uncertainty assumptions. Accordingly, this paper uses Khaled et al. (2018) as the direct compact baseline for instance-matched comparisons, and treats other MILP families as complementary reference paradigms rather than strict one-to-one benchmarks.

#### *2.4. Maintenance constraints*

Maintenance scheduling within aircraft routing has been studied since the early work of Clarke et al. (1997). Aircraft maintenance programs have traditionally grouped tasks into A, B, C, and D check packages with different scopes and frequencies, although operators may use task-based programs rather than these historical labels. The column-generation framework of Barnhart et al. (1998) can embed type-A constraints through constrained shortest-path pricing. Cordeau et al. (2001) apply Benders decomposition to simultaneous routing and crew scheduling with non-linear maintenance resources. Lacasse-Guay et al. (2010) consider maintenance under different business-process models. Section 5 develops one explicit four-level counter-reset formulation; we do not claim that hierarchical maintenance has not appeared in other routing frameworks.

#### *2.5. Heuristic approaches*

Constructive heuristics for TAP are comparatively under-studied relative to exact methods. Grönkvist (2005) provides a comprehensive treatment including local-search improvements. The Dijkstra-based and ACO-based

heuristics studied in the present paper are adapted from a space-time network framework for the Locomotive Assignment Problem.

### 2.6. Position of the present paper

The compact formulation of Khaled et al. (2018) represents the reference compact baseline for exact, cyclic-constraint-free TAP in this study. The present work builds directly upon it, providing: (i) a diagnosis of Constraint (13) and an exact state-based replacement, (ii) a hierarchical maintenance extension, and (iii) an instance-matched benchmark of the implemented formulations and heuristics.

## 3. The Basic Tail Assignment Model

We restate the compact model of Khaled et al. (2018) using the notation of Table 1. This section covers the model without maintenance constraints; they are introduced in Section 4.

### 3.1. Decision variables and objective

Let

$$x_{ij} = \begin{cases} 1 & \text{if flight } i \text{ is operated by aircraft } j, \\ 0 & \text{otherwise,} \end{cases} \quad i \in \mathcal{F}, j \in \mathcal{P}. \quad (1)$$

The objective minimises total assignment cost:

$$\min \sum_{i \in \mathcal{F}} \sum_{j \in \mathcal{P}} c_{ij} x_{ij}. \quad (2)$$

### 3.2. Constraints

*C1 — Flight coverage..* Each flight must be assigned to exactly one aircraft:

$$\sum_{j \in \mathcal{P}} x_{ij} = 1, \quad \forall i \in \mathcal{F}. \quad (3)$$

Table 1: Notation summary.

| Symbol                            | Meaning                                                                             |
|-----------------------------------|-------------------------------------------------------------------------------------|
| $\mathcal{F} = \{1, \dots, n_f\}$ | Set of flight legs                                                                  |
| $\mathcal{P} = \{1, \dots, n_p\}$ | Set of aircraft (tail numbers)                                                      |
| $\mathcal{A} = \{1, \dots, n_k\}$ | Set of airports                                                                     |
| $a_i^\uparrow, a_i^\downarrow$    | Departure / arrival airport of flight $i$                                           |
| $t_i^\uparrow, t_i^\downarrow$    | Departure / arrival time of flight $i$ (minutes)                                    |
| $d_i, \bar{d}_i$                  | Departure / arrival day derived from flight times                                   |
| $\Delta t$                        | Minimum ground turn time (minutes)                                                  |
| $c_{ij}$                          | Assignment cost of flight $i$ to aircraft $j$                                       |
| $a_j^0$                           | Initial position airport of aircraft $j$                                            |
| $F_k^\downarrow(\theta)$          | $\{i \in \mathcal{F} : a_i^\downarrow = k, t_i^\downarrow \leq \theta - \Delta t\}$ |
| $F_k^\uparrow(\theta)$            | $\{i \in \mathcal{F} : a_i^\uparrow = k, t_i^\uparrow < \theta\}$                   |

*C2 — Equipment continuity (non-home airports)..* For any aircraft  $j$ , any non-home airport  $k \neq a_j^0$ , and any flight  $i$  departing from  $k$ , aircraft  $j$  can be assigned to flight  $i$  only if it has already landed at  $k$  early enough:

$$\sum_{i' \in F_k^\downarrow(t_i^\uparrow)} x_{i'j} - \sum_{i' \in F_k^\uparrow(t_i^\uparrow)} x_{i'j} \geq x_{ij}, \quad \forall j \in \mathcal{P}, k \in \mathcal{A} \setminus \{a_j^0\}, i \in \mathcal{F} : a_i^\uparrow = k. \quad (4)$$

*C3 — Equipment continuity (home airport)..* At the home airport  $k = a_j^0$ , one free unit is available at time zero (the initial position):

$$\sum_{i' \in F_k^\downarrow(t_i^\uparrow)} x_{i'j} - \sum_{i' \in F_k^\uparrow(t_i^\uparrow)} x_{i'j} \geq x_{ij} - 1, \quad \forall j \in \mathcal{P}, k = a_j^0, i \in \mathcal{F} : a_i^\uparrow = k. \quad (5)$$

*Integrality..*

$$x_{ij} \in \{0, 1\}, \quad \forall i \in \mathcal{F}, j \in \mathcal{P}. \quad (6)$$

### 3.3. Validity

Constraints (4)–(5) are intended to encode both equipment continuity (C2 of Khaled et al. (2018)) and turn-time feasibility (C4). The set  $F_k^\downarrow(t_i^\uparrow)$  counts only arrivals at least  $\Delta t$  minutes before the departure of flight  $i$ , so a prior arrival can supply flight  $i$  only after the required turn time. This mechanism is valid when departures are strictly ordered, but it does not prevent two equal-time departures from using the same available aircraft, as shown below.

**Proposition 1** (Claimed in Khaled et al. 2018, Proposition 1). *Any feasible solution to (3)–(6) corresponds to a feasible flight plan satisfying conditions C1–C4.*

The proposition is not valid as stated when distinct flights may have equal departure times. The following counterexample identifies the missing case.

### 3.4. Counter-example with simultaneous departures

The balance constraints (4)–(5) use  $F_k^\uparrow(\theta) = \{i \in \mathcal{F} : a_i^\uparrow = k, t_i^\uparrow < \theta\}$  with a strict inequality in departure time. If two flights leave the same airport at the same time, neither flight appears in the other’s departure-balance term. This can make an assignment feasible in (3)–(6) while being physically impossible as a flight plan.

*Example 1* (Feasible to (3)–(6), infeasible in practice). Consider one aircraft  $j$  with initial airport  $a_j^0 = B$  and turn time  $\Delta t = 30$  minutes. Let the three

flights be:

$$\begin{aligned} i_0 : B &\rightarrow A, & t_{i_0}^\uparrow &= 500, & t_{i_0}^\downarrow &= 550, \\ i_1 : A &\rightarrow B, & t_{i_1}^\uparrow &= 600, & t_{i_1}^\downarrow &= 700, \\ i_2 : A &\rightarrow C, & t_{i_2}^\uparrow &= 600, & t_{i_2}^\downarrow &= 700. \end{aligned}$$

Set  $x_{i_0j} = x_{i_1j} = x_{i_2j} = 1$ .

Coverage (3) is satisfied (single aircraft, all flights assigned). Flight  $i_0$  satisfies the home-airport constraint (5), because no flight has arrived at or departed from B before time 500. For  $i_1$  at the non-home airport A:

$$F_A^\downarrow(t_{i_1}^\uparrow) = \{i_0\}, \quad F_A^\uparrow(t_{i_1}^\uparrow) = \emptyset,$$

because  $i_0$  departs from B and  $t_{i_2}^\uparrow = 600 \not\leq 600$ . Hence constraint (4) gives

$$\sum_{i' \in F_A^\downarrow(t_{i_1}^\uparrow)} x_{i'j} - \sum_{i' \in F_A^\uparrow(t_{i_1}^\uparrow)} x_{i'j} = 1 - 0 = 1 \geq x_{i_1j} = 1.$$

Symmetrically for  $i_2$ ,  $t_{i_1}^\uparrow = 600 \not\leq 600$ , so the same calculation gives  $1 \geq x_{i_2j} = 1$ . Therefore (3)–(6) are feasible.

However, the resulting plan is physically infeasible: one aircraft cannot operate both  $i_1$  and  $i_2$  simultaneously at time 600.

*Implication..* In implementations, this corner case is eliminated by adding pairwise non-overlap constraints for time-overlapping flights on the same aircraft,  $x_{i_1j} + x_{i_2j} \leq 1$ .

*Corrected C2–C4 feasibility criterion..* Define  $\mathcal{O}$  as the set of flight pairs that cannot be operated by one aircraft because their operating intervals, including the turn-time buffer, overlap:

$$\mathcal{O} = \left\{ \{i_1, i_2\} \subseteq \mathcal{F} : [t_{i_1}^\uparrow, t_{i_1}^\downarrow + \Delta t) \cap [t_{i_2}^\uparrow, t_{i_2}^\downarrow + \Delta t) \neq \emptyset \right\}. \quad (7)$$

For every pair in this conflict set, impose

$$x_{i_1j} + x_{i_2j} \leq 1, \quad \forall \{i_1, i_2\} \in \mathcal{O}, \forall j \in \mathcal{P}. \quad (8)$$

The corrected core model is therefore (3)–(5), (8), and (6). This preserves the original  $n_f n_p$  binary-variable space while restoring physical executability. It does not, in general, preserve the linear constraint count: the overlap family contains  $|\mathcal{O}|n_p$  inequalities.

### 3.5. Model size

**Proposition 2** (Khaled et al. 2018, Proposition 2). *The original model (2)–(6) contains  $n_f n_p$  binary decision variables and  $O(n_f n_p)$  constraints.*

The corrected model retains  $n_f n_p$  binary variables and has  $O(n_f n_p + |\mathcal{O}|n_p)$  constraints. In the worst case  $|\mathcal{O}| = O(n_f^2)$ , although timetable sparsity normally makes the conflict set much smaller. The original variable count compares favourably with Levin’s  $O(n_f^2 n_p)$ -variable model; Khaled et al. (2018) also report that the time-space network formulation uses approximately 1.25–2 times as many variables and constraints on their test instances.

### 3.6. Illustrative example

*Example 2.* Consider the six-flight, two-aircraft, two-airport instance of Khaled et al. (2018) (Table 2 therein, replicated in Table 2). Turn time  $\Delta t = 30$  min; initial positions airport A for aircraft 1 and airport B for aircraft 2. The optimal assignment (objective = 37,647) is: aircraft 1 serves flights  $1 \rightarrow 2 \rightarrow 3$ ; aircraft 2 serves flights  $4 \rightarrow 5 \rightarrow 6$ . Imposing the cyclic constraint renders this instance infeasible, illustrating the practical advantage of the non-cyclic formulation.

Table 2: Six-flight timetable for Example 2.

| Flight | Dep. airport | Arr. airport | Dep. time | Arr. time |
|--------|--------------|--------------|-----------|-----------|
| 1      | A            | B            | 09:00     | 11:00     |
| 2      | B            | A            | 11:30     | 13:30     |
| 3      | A            | B            | 14:00     | 16:00     |
| 4      | B            | A            | 09:30     | 11:30     |
| 5      | A            | B            | 12:00     | 14:00     |
| 6      | B            | A            | 14:30     | 16:30     |

#### 4. The Maintenance-Extended Model and Correction of Constraint (13)

##### 4.1. Maintenance problem definition

Following Khaled et al. (2018), we consider type-A checks triggered after at most  $T_{\max}^A$  cumulative flight hours. Checks take place at night after the last flight on a given day in an airport with maintenance capacity.

Additional notation is collected in Table 3.

##### 4.2. Additional decision variables

$$z_{ijd} = \begin{cases} 1 & \text{night maintenance at arrival airport of } i, \text{ aircraft } j, \text{ day } d, \\ 0 & \text{else;} \end{cases}$$

$$y_{jd} = \begin{cases} 1 & \text{maintenance of aircraft } j \text{ on day } d, \\ 0 & \text{else,} \end{cases}$$

$$h_{jd} \geq 0$$

accumulated type-A flying

Table 3: Additional notation for the maintenance model.

| Symbol                              | Meaning                                                                             |
|-------------------------------------|-------------------------------------------------------------------------------------|
| $\mathcal{D} = \{1, \dots, n_d\}$   | Planning days                                                                       |
| $d_i$                               | Departure day of flight $i$                                                         |
| $\bar{d}_i$                         | Arrival day of flight $i$                                                           |
| $\mathcal{M} \subseteq \mathcal{A}$ | Maintenance-capable airports                                                        |
| $\text{mcap}_{m,d}$                 | Maintenance capacity of airport $m$ on day $d$                                      |
| $mc_{ijd}$                          | Cost of night maintenance at arrival airport of flight $i$ , aircraft $j$ , day $d$ |
| $F_d^\downarrow(m)$                 | Flights arriving at airport $m$ on day $d$                                          |
| $F(d, i)$                           | $\{i' : d_{i'} = d, t_{i'}^\uparrow > t_i^\downarrow\}$                             |
| $T_{\max}^A$                        | Max cumulative flight minutes before type-A check                                   |
| $d_{\max}$                          | Max calendar-day interval between checks                                            |
| $M_{dd'}$                           | Big- $M$ for day pair $(d, d')$                                                     |
| $F_{d,d'}$                          | Flights with departure day in $(d, d']$                                             |

The variable  $z_{ijd}$  is instantiated only when flight  $i$  arrives at a maintenance-capable airport on day  $d$ . Thus its admissible domain is  $i \in \bigcup_{m \in \mathcal{M}} F_d^\downarrow(m)$ . We assume in this type-A model that a night check fits within the overnight ground period; multi-day checks are treated separately in Section 5.

#### 4.3. Extended model

The objective becomes:

$$\min \sum_{i \in \mathcal{F}} \sum_{j \in \mathcal{P}} c_{ij} x_{ij} + \sum_{d \in \mathcal{D}} \sum_{m \in \mathcal{M}} \sum_{i \in F_d^\downarrow(m)} \sum_{j \in \mathcal{P}} mc_{ijd} z_{ijd}. \quad (9)$$

Constraints (3)–(5) are retained. Additional constraints:

*C8 — Maintenance blocks same-day subsequent flights..*

$$z_{ijd} + x_{i'j} \leq 1, \quad \forall d \in \mathcal{D}, m \in \mathcal{M}, i \in F_d^\downarrow(m), i' \in F(d, i), j \in \mathcal{P}. \quad (10)$$

*C9 — Maintenance requires assignment..*

$$x_{ij} \geq z_{ijd}, \quad \forall d \in \mathcal{D}, m \in \mathcal{M}, i \in F_d^\downarrow(m), j \in \mathcal{P}. \quad (11)$$

*C10 — Capacity..*

$$\sum_{i \in F_d^\downarrow(m)} \sum_{j \in \mathcal{P}} z_{ijd} \leq \text{mcap}_{m,d}, \quad \forall d \in \mathcal{D}, m \in \mathcal{M}. \quad (12)$$

*C11 — Aggregation ( $z \rightarrow y$ )..*

$$\sum_{m \in \mathcal{M}} \sum_{i \in F_d^\downarrow(m)} z_{ijd} = y_{jd}, \quad \forall d \in \mathcal{D}, j \in \mathcal{P}. \quad (13)$$

*C12 — Maximum-spacing constraint..*

$$\sum_{r \in [d, d']} y_{jr} \geq 1, \quad \forall d, d' \in \mathcal{D} \text{ s.t. } d' - d = d_{\max} - 1, \forall j \in \mathcal{P}. \quad (14)$$

*C14 — At most one maintenance event per aircraft per day..*

$$\sum_{m \in \mathcal{M}} \sum_{i \in F_d^\downarrow(m)} z_{ijd} \leq 1, \quad \forall j \in \mathcal{P}, d \in \mathcal{D}. \quad (15)$$

For binary  $y$ , C14 is implied by C11; it is retained because it makes the one-event-per-day rule explicit and remains useful in implementations that relax or disaggregate the aggregation equality.

#### 4.4. Impact of corrected C2–C4 feasibility on C8–C14

The correction introduced in Section 3 (non-overlap constraint (8)) changes the admissible assignment set for  $x$ , and therefore the interpretation of all maintenance constraints layered on top of  $x$ .

*C8..* No algebraic change is required in (10); however, with (8) active, the blocking relation now operates on physically realizable trajectories only. This removes spurious cases where maintenance placement was evaluated on infeasible simultaneous departures.

*C9–C12 and C14..* Constraints (11)–(14), (15) remain structurally unchanged. Their feasible region is tightened indirectly because  $z$  and  $y$  are linked to a strictly smaller, physically valid  $x$ -space.

*C13 and onward..* Hour-accumulation constraints are evaluated over a physically executable  $x$ -space once (8) is imposed. This correction is necessary, but it does not by itself establish that a particular C13 algebraic form correctly propagates the maintenance counter; that property is analysed below.

In summary, the C2–C4 correction primarily changes the domain over which the maintenance constraints operate. It is independent of the separate algebraic correction required for cumulative-hour propagation.

#### *4.5. Rigorous feasibility theorem for C1–C14*

We now state the corrected sufficiency result for the type-A maintenance model. Consider the constraint family (3)–(5), (8), (10)–(14), (17)–(21), (15), and the stated variable domains.

**Theorem 1** (Physical feasibility under corrected C1–C14). *Any solution satisfying the above constraints, with  $x$ ,  $z$ , and  $y$  binary and  $h$  continuous, induces for each aircraft  $j$  a physically executable trajectory over the planning horizon: flights assigned to  $j$  are temporally non-overlapping, satisfy airport continuity and turn time, and admit a day-by-day maintenance schedule satisfying capacity and hour-spacing rules.*

*Proof.* Fix an aircraft  $j$  and let  $S_j = \{i \in \mathcal{F} : x_{ij} = 1\}$  ordered by nondecreasing departure time.

**(i) Flight assignment and continuity.** By (3), every flight is assigned to exactly one aircraft. By (4)–(5), each departure assigned to  $j$  has sufficient prior balance at its departure airport, with the initial unit at  $a_j^0$ . Therefore airport continuity holds along  $S_j$ .

**(ii) Temporal executability.** For any overlapping pair  $\{i_1, i_2\} \in \mathcal{O}$ , (8) enforces  $x_{i_1j} + x_{i_2j} \leq 1$ . Hence no two flights assigned to  $j$  overlap once turn time is included. Combined with (i), this yields a realizable single-aircraft flight timeline.

**(iii) Maintenance event consistency.** By (11), maintenance trigger variables can only be activated on assigned flights. By (13),  $y_{jd} = 1$  iff exactly one daily trigger is selected in aggregate; by (15), at most one maintenance event is attached to aircraft  $j$  on day  $d$ .

**(iv) Maintenance capacity and blocking.** By (12), airport-day capacities are respected. By (10), if maintenance is triggered after flight  $i$  on day  $d$ , then any subsequent same-day flight is forbidden for the same aircraft. Under the stated overnight-duration assumption, the maintenance action is operationally executable.

**(v) Hour/day maintenance admissibility.** By (14), checks cannot be spaced farther than  $d_{\max}$  days. By (17)–(21), the cumulative flight-time state increases by the flying time assigned each day, resets after a maintenance night, and never exceeds  $T_{\max}^A$ .

Steps (i)–(v) show that every aircraft has a consistent, non-overlapping flight sequence with feasible maintenance placement and resource usage.

Therefore the global solution is physically feasible.  $\square$

#### 4.6. Constraint (13): original formulation and its flaw

The original paper's cumulative-flight-time constraint is:

$$\sum_{i \in F_{d,d'}} x_{ij} \tilde{t}_i \leq T_{\max}^A + M_{dd'}(2 - y_{jd} - y_{jd'}) + M_{dd'} \sum_{r \in [d+1, d'-1]} y_{jr}, \quad (13_{\text{paper}})$$

for all  $j \in \mathcal{P}$ , all pairs  $(d, d')$  with  $2 \leq d' - d \leq d_{\max} - 1$ ,  $d < n_d$ , where  $\tilde{t}_i$  is the flight duration of leg  $i$ .

**Lemma 1** (Flaw in  $(13_{\text{paper}})$ ). *Consider two consecutive maintenance events at days  $d$  and  $d'$  with no check between them ( $y_{jr} = 0$  for  $r \in (d, d')$ ). Then the RHS of  $(13_{\text{paper}})$  equals  $T_{\max}^A + M_{dd'}(2 - 1 - 1) + M_{dd'} \cdot 0 = T_{\max}^A$ , which correctly imposes the flight-hour limit between the two checks.*

*However, if  $y_{jd} = 0$  (no maintenance at the start day) and  $y_{jd'} = 1$  (maintenance only at the end day), the RHS becomes  $T_{\max}^A + M_{dd'}(2 - 0 - 1) = T_{\max}^A + M_{dd'}$ , which is vacuously large and imposes no limit on cumulative flight hours before the first check of the period. Hence the constraint fails to enforce the check-interval bound for the opening sub-period.*

*Proof.* Direct substitution of  $y_{jd} = 0$ ,  $y_{jd'} = 1$ ,  $\sum_r y_{jr} = 0$  into  $(13_{\text{paper}})$  gives the stated RHS. Since  $M_{dd'}$  is chosen as the maximum possible cumulative flight time over the horizon (a large positive constant), the resulting constraint is trivially satisfied regardless of the actual flight hours flown.  $\square$

#### 4.7. Endpoint-split strengthening and its limitation

An endpoint-split strengthening considered for replacing  $(13_{\text{paper}})$  consists of two constraints, each relaxed by one endpoint:

$$\sum_{i \in F_{d,d'}} x_{ij} \tilde{t}_i \leq T_{\max}^A + M_{dd'} y_{jd} + M_{dd'} \sum_{r \in [d+1, d'-1]} y_{jr}, \quad (13a)$$

$$\sum_{i \in F_{d,d'}} x_{ij} \tilde{t}_i \leq T_{\max}^A + M_{dd'} y_{jd'} + M_{dd'} \sum_{r \in [d+1, d'-1]} y_{jr}, \quad (13b)$$

for all  $j \in \mathcal{P}$ ,  $(d, d') \in \mathcal{D}^2$  with  $2 \leq d' - d \leq d_{\max} - 1$ ,  $d < n_d$ .

**Lemma 2** (The endpoint split is not sufficient). *With  $F_{d,d'}$  defined as the flights departing in  $(d, d']$ , constraints (13a)–(13b) do not necessarily bound all flying time between two consecutive maintenance nights.*

*Proof.* Let  $T_{\max}^A = 10$ ,  $M_{dd'} = 20$ ,  $d_{\max} \geq 4$ , and consider days  $1, \dots, 4$  with  $y_{j1} = y_{j4} = 1$  and  $y_{j2} = y_{j3} = 0$ . Assign 10 units of flying time on day 2 and 10 units on day 4. For  $(d, d') = (1, 3)$ , the accumulated time is 10 and the second split inequality is active. For  $(2, 4)$ , the accumulated time is 10 and the first split inequality is active. For  $(1, 4)$ , the accumulated time is 20, but both inequalities have right-hand side  $T_{\max}^A + M_{14} = 30$  because both endpoints contain checks. Thus every split inequality is satisfied although 20 units are flown between the consecutive checks, exceeding the limit 10.  $\square$

The original and split formulations are therefore not ordered by strength: the original row is active when both endpoints contain checks, whereas the split rows are active when at least one corresponding endpoint indicator is zero. Neither pattern alone gives an exact cumulative counter.

*Remark 1* (Implementation status). The default builder in `src/model.py` constructs the endpoint-split inequalities (13a)–(13b); the original single-row form (13<sub>paper</sub>) remains available as an explicit, opt-in `use_paper_c13` flag for controlled comparison, and an interval-specific tight big- $M$  (`use_tight_c13_m`)

is available as a formulation-preserving strengthening of either row family. Section 7.4 numerically confirms both Lemma 1 and Lemma 2 against the real built constraint objects of both forms, at two adversarial check patterns on the same four-day test instance: with checks at both window endpoints, all 42 split rows are satisfied while the true cumulative flying time exceeds  $T_{\max}^A$  by a factor of two, and the single paper-form row correctly rejects it; with only the end-day checked, the paper-form row is satisfied at the same violation while all 42 split rows correctly reject it. The event-based implementation in `src/event_model.py` uses a flight-sequence state recursion analogous to the exact formulation below.

*Four-day counterexample and implemented repair..* The failure and its repair can be traced completely on one aircraft over four days. Let  $T_{\max}^A = 10$ ,  $M_{dd'} = 20$ ,  $v_{j1} = v_{j3} = 0$ , and  $v_{j2} = v_{j4} = 10$ , with maintenance after the flights of days 1 and 4:  $y_{j1} = y_{j4} = 1$  and  $y_{j2} = y_{j3} = 0$ . The check on night 1 resets the counter before day 2, whereas the check on night 4 occurs only after the day-4 flying and therefore cannot legalize those hours. Nevertheless, every endpoint-split row is satisfied. For  $(d, d') = (1, 3)$ , the interval  $(1, 3]$  contains 10 units and the row anchored at day 3 has right-hand side 10. For  $(2, 4)$ , the interval  $(2, 4]$  also contains 10 units and the row anchored at day 2 has right-hand side 10. For  $(1, 4)$ , the interval contains all 20 units, but each row has right-hand side  $10 + 20 = 30$  because its corresponding endpoint contains a check. Thus the split model accepts 20 units between consecutive maintenance nights. Under the exact recursion (17)–(21), the same data give  $h_{j2} = 10$ ,  $h_{j3} = 10$ , and  $h_{j4} = 20$ ; the last value violates  $h_{j4} \leq T_{\max}^A = 10$ . In the code, the preserved class `LegacyEndpointSplitMILPScheduler` is identi-

ified as `legacy_endpoint_split`. The separate `EventMILPScheduler`, identified as `event_exact_state`, implements the equivalent event-level repair: its state is initialized with the pre-horizon hours, propagated by each selected successor flight, reset to the next flight's duration only when a qualifying maintenance event follows the preceding flight, and bounded by the applicable A/B threshold at every event. Unselected successor arcs are deactivated with the valid bound  $T_{\max}^c + \tilde{t}_\ell$ , while an unused source arc is deactivated with  $\Phi_j^c + \tilde{t}_i$ ; using only  $T_{\max}^c$  in either role can create false infeasibility. The comparison harness `experiments/compare_formulations.py` selects these two identifiers explicitly and rejects unknown identifiers, so the legacy and exact-state rows can no longer be produced by the same builder.

#### 4.8. Exact cumulative-state formulation

Define the flying time assigned to aircraft  $j$  on day  $d$  by

$$v_{jd} = \sum_{i \in \mathcal{F}: d_i = d} \tilde{t}_i x_{ij}. \quad (16)$$

Let  $h_{jd} \geq 0$  be the accumulated type-A flying time immediately after the flights of day  $d$  and before maintenance that night. Set  $h_{j0} = \Phi_j^A$  and  $y_{j0} = 0$ . For consecutive days, impose

$$h_{jd} \geq h_{j,d-1} + v_{jd} - M^H y_{j,d-1}, \quad (17)$$

$$h_{jd} \leq h_{j,d-1} + v_{jd} + M^H y_{j,d-1}, \quad (18)$$

$$h_{jd} \geq v_{jd}, \quad (19)$$

$$h_{jd} \leq v_{jd} + M^H (1 - y_{j,d-1}), \quad (20)$$

$$0 \leq h_{jd} \leq T_{\max}^A, \quad \forall j \in \mathcal{P}, d \in \mathcal{D}, \quad (21)$$

where  $M^H \geq T_{\max}^A$  is valid because  $h_{j,d-1} \leq T_{\max}^A$ . If  $y_{j,d-1} = 0$ , the first two inequalities give  $h_{jd} = h_{j,d-1} + v_{jd}$ . If  $y_{j,d-1} = 1$ , the reset inequalities give  $h_{jd} = v_{jd}$ . The upper bound then enforces the flight-hour threshold on every day, including the flights operated immediately before a night check.

#### 4.9. Pre-horizon accumulated hours

The original paper does not account for flying time accumulated before the planning horizon. In the cumulative-state formulation this history is included exactly through the initial condition

$$h_{j0} = \Phi_j^A, \quad 0 \leq \Phi_j^A \leq T_{\max}^A, \quad \forall j \in \mathcal{P}. \quad (22)$$

The condition is required for every aircraft, including  $\Phi_j^A = 0$ ; no separate opening-window big- $M$  inequality is needed.

#### 4.10. Integrality constraints

$$\begin{aligned} x_{ij} \in \{0, 1\}, \quad z_{ijd} \in \{0, 1\}, \quad y_{jd} \in \{0, 1\}, \quad (i, j) \in \mathcal{F} \times \mathcal{P}, \\ d \in \mathcal{D}, \quad i \in \bigcup_{m \in \mathcal{M}} F_d^\downarrow(m) \text{ for } z_{ijd}. \end{aligned} \quad (23)$$

The variables  $h_{jd}$  are continuous and satisfy  $0 \leq h_{jd} \leq T_{\max}^A$  as imposed in (21).

### 5. Extension to the Full $\{A, B, C, D\}$ Maintenance Hierarchy

We model a traditional packaged maintenance program with four check levels. Type-A and B checks are triggered by cumulative flight hours; type-C

and D are triggered by calendar days. A heavier check subsumes all lighter checks (a D-check also satisfies the C-, B-, and A-check requirement for the same period). This hierarchy is a modeling convention for the benchmark, not a claim that regulators require every operator to use these labels or triggers.

### 5.1. Extended notation

We introduce the check-type index set  $\mathcal{C} = \{A, B, C, D\}$  and the binary variable

$$y_{jdc} = \begin{cases} 1 & \text{aircraft } j \text{ undergoes check } c \text{ on day } d, \\ 0 & \text{else.} \end{cases} \quad (24)$$

A *hierarchy indicator* variable

$$\eta_{jdc} = \begin{cases} 1 & \text{aircraft } j \text{ undergoes a check of level } \geq c \text{ on day } d, \\ 0 & \text{else,} \end{cases} \quad (25)$$

is defined such that  $\eta_{jdc} = 1$  whenever  $y_{jdc'} = 1$  for any  $c' \succeq c$  (where  $D \succ C \succ B \succ A$ ).

### 5.2. Hierarchy constraints

*H1 — Heaviest check subsumes lighter checks..*

$$\eta_{jdc} \geq \eta_{jdc'}, \quad \forall j, d, c \prec c'. \quad (26)$$

Thus, for example, a type-D event makes all four effective-check indicators equal to one, while a type-A event affects only  $\eta_{jdA}$ .

*H2 — Check indicator consistency..*

$$\eta_{jdc} = \sum_{c' \succeq c} y_{jdc'}, \quad \forall j, d, c \quad (\text{exactly one check level per event, or none}). \quad (27)$$

In practice implemented as:  $\eta_{jdA} = y_{jdA} + y_{jdB} + y_{jdC} + y_{jdD}$ ,  $\eta_{jdB} = y_{jdB} + y_{jdC} + y_{jdD}$ , etc.

For multi-day checks, let  $s_{jmdc} = 1$  when check  $c \in \{C, D\}$  starts for aircraft  $j$  at maintenance airport  $m$  on day  $d$ , and impose  $y_{jdc} = \sum_{m \in \mathcal{M}} s_{jmdc}$ . Let  $q_{jmrc} = 1$  when that aircraft occupies the maintenance resource at airport  $m$  on day  $r$ . Both  $s$  and  $q$  are binary.

*H3 — At most one check type per aircraft per day..*

$$\sum_{c \in \mathcal{C}} y_{jdc} \leq 1, \quad \forall j \in \mathcal{P}, d \in \mathcal{D}. \quad (28)$$

### 5.3. Flight-hour triggered checks (A and B)

For each  $c \in \{A, B\}$ , let  $h_{jd}^c$  be the accumulated flying time after day  $d$  and before that night's maintenance, let  $h_{j0}^c = \Phi_j^c$ , and use the daily flying time  $v_{jd}$  from (16). The exact recursion is

$$h_{jd}^c \geq h_{j,d-1}^c + v_{jd} - M^c \eta_{j,d-1,c}, \quad (29)$$

$$h_{jd}^c \leq h_{j,d-1}^c + v_{jd} + M^c \eta_{j,d-1,c}, \quad (30)$$

$$h_{jd}^c \geq v_{jd}, \quad (31)$$

$$h_{jd}^c \leq v_{jd} + M^c(1 - \eta_{j,d-1,c}), \quad (32)$$

$$0 \leq h_{jd}^c \leq T_{\max}^c, \quad \forall j \in \mathcal{P}, d \in \mathcal{D}, c \in \{A, B\}. \quad (33)$$

A type-B, C, or D event resets the A counter because it activates  $\eta_{jdA}$ ; type-C and D events similarly reset both A and B counters. No calendar-day spacing

constraint is imposed on A/B checks unless the maintenance policy supplies a separate calendar threshold in addition to the flight-hour threshold.

#### 5.4. Calendar-day triggered checks (C and D)

For checks  $c \in \{C, D\}$ , the constraint is on calendar-day spacing only:

*Maximum-spacing (C12 for C/D)..*

$$\sum_{r \in [d, d']} \eta_{jrc} \geq 1, \quad d' - d = d_{\max}^c - 1, \quad \forall j, c \in \{C, D\}. \quad (34)$$

*Multi-day check duration (C14b)..* A C- or D-check occupies multiple consecutive days. Let  $\delta^c$  denote the duration in days. Start and occupancy indicators are linked by

$$q_{jmrc} = \sum_{d \in \mathcal{D}: d \leq r < d + \delta^c} s_{jmdc}, \quad \forall j, m \in \mathcal{M}, r \in \mathcal{D}, c \in \{C, D\}. \quad (35)$$

where starts extending beyond the planning horizon are disallowed unless post-horizon occupancy is explicitly represented. The following rows prevent overlap for one aircraft and enforce airport-day capacity  $\kappa_{mr}$ :

$$\sum_{m \in \mathcal{M}} \sum_{c \in \{C, D\}} q_{jmrc} \leq 1, \quad \forall j, r \in \mathcal{D}, \quad (36)$$

$$\sum_{j \in \mathcal{P}} \sum_{c \in \{C, D\}} q_{jmrc} \leq \kappa_{mr}, \quad \forall m \in \mathcal{M}, r \in \mathcal{D}. \quad (37)$$

The start variable  $s_{jmdc}$  is permitted only when the aircraft can reach  $m$  by day  $d$ ; this uses the same assignment-to-location linking principle as C9. Capacity is imposed on occupancy, not only on starts.

Under the corrected basic feasibility criterion (Equation (8)), (35) is interpreted over physically executable aircraft trajectories; no additional linearization is needed.

*C15 — No flight during active maintenance..* No flight operating while a C- or D-check is active may be assigned to the aircraft. Let  $\mathcal{D}(i)$  be the set of calendar days intersected by the operating interval of flight  $i$ ; then

$$x_{ij} \leq 1 - \sum_{m \in \mathcal{M}} q_{jmrc}, \quad \forall i \in \mathcal{F}, j \in \mathcal{P}, r \in \mathcal{D}(i), c \in \{C, D\}. \quad (38)$$

### 5.5. Pre-horizon accumulated values

For aircraft  $j$  and check  $c$ :

- $\Phi_j^c$  ( $c \in \{A, B\}$ ): accumulated flight minutes before horizon start.
- $\Psi_j^c$  ( $c \in \{C, D\}$ ): elapsed calendar days before horizon start.

The values  $\Phi_j^c$  are the initial conditions in (29)–(33). For C/D, the opening sliding windows in (34) are shortened according to  $\Psi_j^c$  so that a check due near the horizon start cannot be deferred by a full new interval.

### 5.6. Model size impact

**Proposition 3.** *Adding the full  $\{A, B, C, D\}$  hierarchy over the basic maintenance model of Section 4 adds  $O(n_p n_d |\mathcal{C}| + n_p n_m n_d |\{C, D\}|)$  check, hierarchy, location, occupancy, and state variables and  $O(n_p n_d |\mathcal{C}| + n_p n_m n_d |\{C, D\}| + n_f n_p |\{C, D\}|)$  linking, capacity, and flight-exclusion constraints. The assignment variables remain unchanged.*

The actual size multiplier depends on the number of feasible maintenance events, the durations  $\delta^c$ , and whether sparse flight-exclusion rows are generated. The current result artifacts do not support a general empirical multiplier for the exact hierarchy formulation.

## 6. Event-Based Reformulation: Removing the Day Index

Sections 4 and 5 keep the day-indexed structure of Khaled et al. (2018): maintenance state is tracked per calendar day  $d \in \mathcal{D}$ , and both the flight-hour recursion (17)–(21) and the hierarchy extension repeat this index for every check type. This section presents an updated, day-free reformulation of the same C1–C14 requirements. It is the model actually implemented as `event_exact_state` in `src/event_model.py` and referenced throughout Sections 4 and 7 as the exact-state formulation; here we give it a complete, self-contained derivation and prove its correctness directly, rather than only by analogy to the day-indexed recursion.

### 6.1. Modelling principle

Every flight already carries an absolute departure and arrival timestamp measured from the start of the planning horizon, so its calendar day is derived data,  $d_i = \lfloor t_i^\uparrow / 1440 \rfloor$ , not a decision index. Once this is recognised, the day set  $\mathcal{D}$  can be dropped from every maintenance variable without losing any information: maintenance is represented as a virtual flight from an airport to itself, starting immediately after a real flight, occupying the aircraft for a fixed duration, and consuming station capacity for that duration. A single binary variable  $z_{ijk}$  selects check type  $k$  after flight  $i$  for aircraft  $j$ ; there is no longer one variable per (flight, aircraft, *day*, check) tuple, and no separate variable is needed for a multi-day check’s later occupancy days.

The implementation uses sparse indexes for these event domains. Successor flights are grouped by origin airport and sorted by departure time, so each connection search begins at the first timestamp satisfying the turn-time

bound. Calendar-check events are likewise indexed by aircraft, check type, and arrival timestamp before C14 rows are generated. These indexes change construction time only; they preserve the same connection and deadline sets and therefore do not alter the formulation or its correctness argument.

## 6.2. Additional sets and parameters

Let  $\mathcal{C}_H = \{A, B\}$  (flight-hour triggered) and  $\mathcal{C}_T = \{C, D\}$  (calendar-time triggered) partition  $\mathcal{C}$ , matching the split already used in Sections 4 and 5. For flight  $i$ , write  $o_i = a_i^\downarrow$ ,  $a_i = a_i^\uparrow$ , and let  $p_i = t_i^\downarrow - t_i^\uparrow$  be its flying time (previously written  $\tilde{t}_i$  in Section 4). For aircraft  $j$ , let  $a_j^0$  be its initial airport and  $r_j$  its initial availability time. Let  $\Delta t$  be the minimum turn time (as in Section 3) and  $\delta_k$  the fixed duration of check  $k$ . Maintenance attached to flight  $i$  starts at  $\sigma_i = t_i^\downarrow$  and is the virtual leg

$$m(i, j, k) = (a_i, a_i, \sigma_i, \sigma_i + \delta_k). \quad (39)$$

Only flights in  $\mathcal{F}^M = \{i \in \mathcal{F} : a_i \in \mathcal{M}\}$  may be followed by maintenance. The sparse connection set generalises the arcs implicit in (4)–(5):

$$\mathcal{E} = \{(i, \ell) \in \mathcal{F} \times \mathcal{F} : a_i = o_\ell, t_\ell^\uparrow \geq t_i^\downarrow + \Delta t\}. \quad (40)$$

For each check type, the connections that leave enough time for that maintenance event are

$$\mathcal{E}_k = \{(i, \ell) \in \mathcal{E} : t_\ell^\uparrow \geq \sigma_i + \delta_k + \Delta t\}. \quad (41)$$

The aircraft-specific set of feasible first flights is

$$\mathcal{F}_j^0 = \{i \in \mathcal{F} : o_i = a_j^0, t_i^\uparrow \geq r_j\}. \quad (42)$$

Let  $T_k$  be the maximum flying time between qualifying checks for  $k \in \mathcal{C}_H$  (equal to  $T_{\max}^c$  of Section 4), and let  $L_k$  be the maximum elapsed time between qualifying checks for  $k \in \mathcal{C}_T$ , so a limit of  $d_k$  calendar days is stored as  $L_k = 1440 d_k$  minutes. Parameters  $\Phi_{jk}$  and  $\Psi_{jk}$  are, respectively, the pre-horizon accumulated flying time and elapsed calendar time (the event-based counterparts of  $\Phi_j^c$  and  $\Psi_j^c$  from Section 4.9). The relation  $k' \succeq k$  means check  $k'$  is at least as heavy as  $k$ , exactly as the hierarchy order of Section 5.

### 6.3. Decision variables and objective

$$x_{ij} \in \{0, 1\}, \quad \text{flight } i \text{ is assigned to aircraft } j \text{ (as in Section 3)}, \quad (43)$$

$$w_{i\ell j} \in \{0, 1\}, \quad \text{aircraft } j \text{ operates } \ell \text{ immediately after } i, \quad (44)$$

$$z_{ijk} \in \{0, 1\}, \quad \text{aircraft } j \text{ receives check } k \text{ immediately after } i, \quad (45)$$

$$u_{ijk} \geq 0, \quad \text{flying time accumulated after } i, \text{ before its check, } k \in \mathcal{C}_H. \quad (46)$$

Variable  $z_{ijk}$  generalises the day-indexed trigger  $z_{ijd}$  of Section 4 by replacing the day index  $d$  with the triggering flight  $i$  itself – the day is recoverable as  $\lfloor \sigma_i/1440 \rfloor$  whenever it is needed for reporting. Likewise  $u_{ijk}$  generalises the per-day state  $h_{jd}^c$  of (29), indexed by the flight after which the state is measured instead of by calendar day. Boundary variables  $w_{0ij}$  and  $w_{i\omega j}$  indicate, respectively, the first and last flight of aircraft  $j$ 's route. The hierarchy-reset indicator is not stored as a variable; it is the derived expression

$$\rho_{ijk} = \sum_{k' \in \mathcal{C}: k' \succeq k} z_{ijk'}, \quad (47)$$

the event-indexed analogue of the hierarchy indicator  $\eta_{jdc}$  of Section 5.

With assignment costs  $c_{ij}$  (as in (2)) and maintenance costs  $g_{ijk}$ , the objective is

$$\min \sum_{i \in \mathcal{F}} \sum_{j \in \mathcal{P}} c_{ij} x_{ij} + \sum_{i \in \mathcal{F}^M} \sum_{j \in \mathcal{P}} \sum_{k \in \mathcal{C}} g_{ijk} z_{ijk}. \quad (48)$$

#### 6.4. Event-based constraints C1–C14

*C1 – Flight coverage..* Identical to (3):

$$\sum_{j \in \mathcal{P}} x_{ij} = 1, \quad \forall i \in \mathcal{F}. \quad (C1)$$

*C2 – Unique predecessor..* Every assigned flight has one predecessor or is the first route flight:

$$\mathbb{1}_{\{\ell \in \mathcal{F}_j^0\}} w_{0\ell j} + \sum_{i: (i, \ell) \in \mathcal{E}} w_{i\ell j} = x_{\ell j}, \quad \forall \ell \in \mathcal{F}, j \in \mathcal{P}. \quad (C2)$$

Source variables  $w_{0\ell j}$  are instantiated only for  $\ell \in \mathcal{F}_j^0$ ; the indicator is notation for this sparse domain, not an extra variable.

*C3 – Unique successor..* Every assigned flight has one successor or is the last route flight:

$$w_{i\omega j} + \sum_{\ell: (i, \ell) \in \mathcal{E}} w_{i\ell j} = x_{ij}, \quad \forall i \in \mathcal{F}, j \in \mathcal{P}. \quad (C3)$$

*C4 – One time-ordered route per aircraft..*

$$\begin{aligned} \sum_{i \in \mathcal{F}} w_{0ij} &\leq 1, & \forall j \in \mathcal{P}, \\ \sum_{i \in \mathcal{F}} w_{i\omega j} &= \sum_{i \in \mathcal{F}} w_{0ij}, & \forall j \in \mathcal{P}. \end{aligned} \quad (C4)$$

Airport continuity, turn time, and the non-overlap correction of Section 3 are all embedded in  $\mathcal{E}$ : an arc  $(i, \ell)$  exists only if  $\ell$  departs from  $i$ 's arrival

airport no earlier than  $t_i^\downarrow + \Delta t$ , which already excludes the simultaneous-departure counterexample of Example 1. Since every arc moves forward in time, C2–C4 cannot produce a subtour, so the cyclic-rotation constraint of the classical time-space formulations is unnecessary here, as in the base model of Section 3.

*C5 – Maintenance requires assignment..*

$$z_{ijk} \leq x_{ij}, \quad \forall i \in \mathcal{F}^M, j \in \mathcal{P}, k \in \mathcal{C}. \quad (\text{C5})$$

*C6 – At most one physical check after a flight..*

$$\sum_{k \in \mathcal{C}} z_{ijk} \leq 1, \quad \forall i \in \mathcal{F}^M, j \in \mathcal{P}. \quad (\text{C6})$$

A heavy check satisfies lighter requirements through  $\rho_{ijk}$  (Equation (47)), so simultaneous A, B, C, and D checks after the same flight are unnecessary, matching H1–H3 of Section 5.

*C7 – Maintenance-airport eligibility..* Variables  $z_{ijk}$  are instantiated only for  $i \in \mathcal{F}^M$ ; a dense equivalent is

$$z_{ijk} = 0, \quad \forall i \in \mathcal{F} \setminus \mathcal{F}^M, j \in \mathcal{P}, k \in \mathcal{C}. \quad (\text{C7})$$

*C8 – Maintenance duration blocks the next flight..* An insufficient connection and a maintenance action cannot both be selected:

$$w_{i\ell j} + z_{ijk} \leq 1, \quad \forall (i, \ell) \in \mathcal{E} \setminus \mathcal{E}_k, j \in \mathcal{P}, k \in \mathcal{C}. \quad (\text{C8})$$

For a terminal flight,  $z_{ijk}$  is instantiated only if  $\sigma_i + \delta_k$  falls inside the planning horizon. A multi-day check is therefore one fixed-duration virtual leg, not one binary variable per occupied day as in the start/occupancy pair  $(s_{jmdc}, q_{jmrc})$  of (35).

*C9 – Continuous-time station capacity..* Let  $b_a(\theta)$  be the number of available maintenance positions at airport  $a$  and time  $\theta$ , and let  $\Theta$  collect all maintenance start times together with every breakpoint at which  $b_a(\theta)$  changes – the only times at which the left- or right-hand side below can change value:

$$\sum_{\substack{i \in \mathcal{F}^M : a_i = a \\ k \in \mathcal{C} : \sigma_i \leq \theta < \sigma_i + \delta_k}} \sum_{j \in \mathcal{P}} z_{ijk} \leq b_a(\theta), \quad \forall a \in \mathcal{M}, \theta \in \Theta. \quad (\text{C9})$$

This replaces the day-granular capacity constraint (12) with an exact continuous-time version; setting  $b_a(\theta) = \text{mcap}_{a,d}$  for  $\theta$  inside day  $d$  recovers (12) as a special case.

*C10 – Maintenance hierarchy..*

$$\rho_{ijk} = \sum_{k' \succeq k} z_{ijk'}, \quad \forall i \in \mathcal{F}^M, j \in \mathcal{P}, k \in \mathcal{C}. \quad (\text{C10})$$

This is the event-indexed restatement of H2 (27); it is a derived expression, not an extra decision variable, and replaces the day-indexed hierarchy indicator  $\eta_{jdc}$ .

*C11 – Flight-hour counter bounds..* For  $k \in \mathcal{C}_H$ , the counter is active only for an assigned flight and never exceeds the threshold:

$$p_i x_{ij} \leq u_{ijk} \leq T_k x_{ij}, \quad \forall i \in \mathcal{F}, j \in \mathcal{P}, k \in \mathcal{C}_H. \quad (\text{C11})$$

*C12 – Flight-hour accumulation and reset..* For each route connection, let  $R_{\ell k} = T_k + p_\ell$  deactivate propagation when the successor arc is not selected; the reset term needs only the state bound  $T_k$ :

$$\begin{aligned} u_{\ell jk} &\geq u_{ijk} + p_\ell - R_{\ell k}(1 - w_{i\ell j}) - T_k \rho_{ijk}, \\ u_{\ell jk} &\leq u_{ijk} + p_\ell + R_{\ell k}(1 - w_{i\ell j}) + T_k \rho_{ijk}, \\ u_{\ell jk} &\leq p_\ell + T_k(2 - w_{i\ell j} - \rho_{ijk}), \quad \forall (i, \ell) \in \mathcal{E}, j \in \mathcal{P}, k \in \mathcal{C}_H. \end{aligned} \quad (\text{C12})$$

If  $w_{i\ell j} = 1$  and no qualifying check follows  $i$  ( $\rho_{ijk} = 0$ ), the first two rows give  $u_{\ell jk} = u_{ijk} + p_\ell$ : pure accumulation. If a qualifying check follows  $i$  ( $\rho_{ijk} = 1$ ), the third row together with C11 gives  $u_{\ell jk} = p_\ell$ : an exact reset. C12 is therefore the event-indexed form of the exact recursion (17)–(21) of Section 4.8, propagated along the aircraft's actual selected route instead of along a fixed calendar-day grid, with no day-pair windows and no endpoint ambiguity – it therefore cannot exhibit the split-form failure mode proved in Lemma 2 and confirmed numerically in Section 7.4, since there C12 has no "window" that both endpoints can simultaneously relax: every successor arc is checked individually.

*C13 – Pre-horizon accumulated flying time..* For a first-flight arc, the counter is initialised from known aircraft history, generalising (22):

$$\begin{aligned} u_{ijk} &\geq \Phi_{jk} + p_i - (\Phi_{jk} + p_i)(1 - w_{0ij}), \\ u_{ijk} &\leq \Phi_{jk} + p_i + (\Phi_{jk} + p_i)(1 - w_{0ij}), \quad \forall i \in \mathcal{F}, j \in \mathcal{P}, k \in \mathcal{C}_H. \end{aligned} \quad (\text{C13})$$

C11–C13 jointly replace the cumulative-flight-hour family (13<sub>paper</sub>)–(22) in full, not only its corrected form.

*C14 – Calendar-time spacing without a day set..* Let  $H_0$  and  $H_1$  be the horizon boundaries. For  $k \in \mathcal{C}_T$ , the first deadline is  $d_{jk}^0 = H_0 + L_k - \Psi_{jk}$ ; if  $d_{jk}^0 \leq H_1$ , require

$$\sum_{i \in \mathcal{F}^M: H_0 \leq \sigma_i \leq d_{jk}^0} \rho_{ijk} \geq 1, \quad \forall j \in \mathcal{P}, k \in \mathcal{C}_T. \quad (\text{C14a})$$

Every selected check whose next deadline lies inside the horizon must be followed by another qualifying check within  $L_k$  time units:

$$\sum_{\ell \in \mathcal{F}^M: \sigma_i < \sigma_\ell \leq \sigma_i + L_k} \rho_{\ell jk} \geq \rho_{ijk}, \quad \forall i \in \mathcal{F}^M: \sigma_i + L_k \leq H_1, j \in \mathcal{P}, k \in \mathcal{C}_T. \quad (\text{C14b})$$

These timestamp windows are the continuous-time counterpart of the sliding day window (34), and give the same due-by-deadline semantics while the virtual leg (39) still blocks the aircraft for the check's full fixed duration.

### 6.5. Variables and constraints made obsolete

The reformulation removes the explicit day index  $d$  and every day-indexed maintenance variable of Sections 4 and 5: the trigger  $z_{ijd}$  (now  $z_{ijk}$ ), the daily indicator  $y_{jdc}$ , the hierarchy indicator  $\eta_{jdc}$  (now the derived  $\rho_{ijk}$ ), and the multi-day start/occupancy pair  $s_{jmdc}/q_{jmrc}$ . It also removes one binary per occupied day of a multi-day check and every day-pair accumulation constraint (13a)–(13b). Route information is now explicit through the sparse successor arcs  $w_{i\ell j}$ , which is what lets C12 propagate the flight-hour counter exactly along the aircraft's actual selected sequence.

The implementation generates these arcs from pre-indexed airport, aircraft, and timestamp lists rather than repeated scans of all flight pairs. The same pre-indexing is used for reachable calendar-check events in C14. The indexes change construction time only: they preserve the connection and deadline sets, so the  $w$  variables remain explicit and the correctness argument is unchanged.

### 6.6. Complexity and memory reduction

Let  $n_f = |\mathcal{F}|$ ,  $n_p = |\mathcal{P}|$ ,  $n_c = |\mathcal{C}|$ ,  $n_d$  be the day-indexed model's number of planning days, and  $n_e = |\mathcal{E}|$ . The day-indexed maintenance tensor of Section 5 contains  $O(n_p n_d n_c)$  check/hierarchy binaries plus  $O(n_p n_m n_d |\mathcal{C}_T|)$  location/occupancy binaries. The event-based tensor contains only  $O(n_f n_p n_c)$  binaries, instantiated sparsely over  $\mathcal{F}^M$ :

$$O(n_p n_d n_c + n_p n_m n_d |\mathcal{C}_T|) \longrightarrow O(n_f n_p n_c). \quad (49)$$

The complete event-based model contains

$$O(n_f n_p + n_e n_p + n_f n_p n_c) \quad (50)$$

binary variables, plus  $O(n_f n_p |\mathcal{C}_H|)$  continuous counters. The sparse successor variables add  $O(n_e n_p)$ , but  $n_e \ll n_f^2$  in practice because airport and time incompatibilities eliminate most arcs. Crucially, the maintenance model no longer scales multiplicatively with the number of planning days  $n_d$ , which is the main driver of model size in Section 5 on long horizons. This variable-count advantage is confirmed empirically in Section 7.6, but with an important caveat: constraint (C12) is itself built over every arc in  $\mathcal{E}$ , so on dense timetables where  $|\mathcal{E}|$  grows quickly the event-based model's *constraint* count can exceed the day-indexed model's, even while its variable count remains smaller – a genuine practical trade-off, not a strict improvement in every dimension.

*Remark 2.* Removing the day dimension does not remove calendar information: every timestamp must remain on one absolute, horizon-wide time scale, or overnight gaps and the C/D elapsed-time requirements of C14 cannot be recovered. This is why the implementation, `EventMILPScheduler`

Table 4: Principal variable counts: day-indexed hierarchy model (Section 5) versus the event-based reformulation of this section.

| Component           | Day-indexed model             | Event-based model                            |
|---------------------|-------------------------------|----------------------------------------------|
| Flight assignment   | $n_f n_p$                     | $n_f n_p$                                    |
| Maintenance choice  | $n_p n_d n_c$                 | $n_f n_p n_c$ (sparse over $\mathcal{F}^M$ ) |
| Multi-day occupancy | $n_p n_m n_d  \mathcal{C}_T $ | removed (virtual-leg duration)               |
| Hierarchy indicator | $n_p n_d n_c$                 | derived, not stored                          |
| Route sequence      | implicit                      | $n_e n_p$ sparse binaries                    |
| Hour-check state    | day-pair rows                 | $2n_f n_p$ continuous variables              |

in `src/event_model.py`, keeps the legacy JSON instance format unchanged and derives  $o_i, a_i, t_i^\uparrow, t_i^\downarrow$  directly from it.

### 6.7. Correctness

**Proposition 4.** *Assume  $\mathcal{E}$  contains exactly the physically feasible ordinary connections (40) and C8 is imposed for every maintenance-incompatible connection. Every integral solution of C1–C14 above then defines, for each aircraft, a time-feasible route whose selected maintenance events satisfy duration, station capacity, hierarchy, cumulative flight-hour limits, and calendar-time limits.*

*Proof.* C1 assigns every flight exactly once. C2–C4 chain assigned flights into a single time-ordered route per aircraft, with turn-time and non-overlap feasibility guaranteed by the construction of  $\mathcal{E}$  (Equation (40)), so no separate non-overlap family analogous to (8) is required here. C5–C8 attach each maintenance event only to an assigned flight at an eligible station and

reserve its full duration by blocking every connection that does not leave enough time. C9 enforces continuous-time station capacity. C10 applies the hierarchy reset exactly as H1–H3 of Section 5. C11–C13 propagate and bound the flight-hour counter along the aircraft’s actual selected route, by the same case analysis as Section 4.8: whenever  $w_{\ell j} = 1$ , exactly one of "pure accumulation" or "exact reset" applies at every successor arc, so the counter can never silently skip a window the way the day-indexed split form of Lemma 2 can. C14 links every qualifying calendar-time check to its predecessor within the required elapsed time, using the known pre-horizon state  $\Psi_{jk}$ . All stated operating and maintenance requirements therefore hold for every aircraft, hence for the whole solution.  $\square$

This completes, with a full independent derivation, the exact-state formulation that Sections 4.6 and 4.7 showed was needed to repair Constraint (13): Proposition 4 is the event-based counterpart of Theorem 1, extended from the single type-A check of Section 4 to the full  $\{A, B, C, D\}$  hierarchy of Section 5, without a day index anywhere in the maintenance constraints.

## 7. Computational Study

### 7.1. Implementation and environment

All models are implemented in Python 3.10+ using the Pyomo optimisation framework (Hart et al., 2023). The compact benchmark in Table 5 and the matched diagnostic in Table 10 were solved with IBM ILOG CPLEX (IBM, 2023) through Pyomo’s `cplex` shell backend. For the new Constraint (13) verification and per-check-type hierarchy study (Sections 7.4 and 7.5) results are cross-validated with two independent exact MILP solvers:

the open-source HiGHS solver (Huangfu and Hall, 2018) through Pyomo’s `appsi_highs` interface, and the full, unrestricted IBM ILOG CPLEX 12.6 build distributed with AMPL through Pyomo’s `cplexamp` ASL interface. The CPLEX copy already on this machine’s PATH is a size-restricted Community/Preview edition (1 000 variable/constraint limit) that raises **CPLEX Error 1016** on every instance in this study once maintenance constraints are included; it is not used for any reported result. HiGHS and the unrestricted CPLEX build agree exactly (same status, same objective) on every instance reported in Section 7.5, giving solver-independent confidence in the correctness claims. The code, instance generator, and experiment runner are publicly available at the companion repository.

*Comparison protocol.* The compact benchmark and the matched diagnostic use the same JSON instances and assignment-cost data for both formulations where a comparison is claimed. They do not by themselves certify cumulative-hour correctness; Sections 7.4 and 7.5 evaluates the real built constraints and independently replays maintenance counters.

## 7.2. Instance generation

The current generator creates deterministic synthetic instances. For parameters  $(\delta, n_p, n_d)$  it generates approximately  $\delta n_p n_d$  flights, chooses airports and times from seeded pseudo-random distributions, and uses a stable seed for each test index. Thus  $\delta$  is a flight-count scaling parameter in this repository, not the fleet-to-minimum- fleet ratio used by Khaled et al. (2018). Unless a table states otherwise, the quick snapshots contain one generated instance per parameter cell.

Table 5: Measured MILP benchmark on `ABCD_capacity_bottleneck_test`. This compact row gives the actual solver statistics for one exact solve.

| Status  | CPU (s) | Wall (s) | Obj   |
|---------|---------|----------|-------|
| optimal | 0.09    | 0.20     | 8 300 |

### 7.3. Reference benchmark

For a compact numerical anchor, Table 5 reports one exact solve on the small `ABCD_capacity_bottleneck_test` instance. The former random snapshots that reproduced the configurations of Tables 6, 10, and 11 of Khaled et al. (2018) are omitted: their infeasible cells were neither matched across formulations nor informative for the correctness claim.

### 7.4. Numeric verification of the Constraint (13) endpoint-split flaw

Lemmas 1 and 2 are proved algebraically in Sections 4.6 and 4.7. We additionally verify both directly against the production implementation: `experiments/c13_loophole_v` builds the real `LegacyEndpointSplitMILPScheduler` model – with both the split formulation and, using the restored `use_paper_c13` toggle, the original Khaled et al. (2018) form – on the four-day, single-aircraft instance `data/instances/c13_loophole_test.json` ( $T_{\max}^A = 600$  min, 600 min flown on each of days 2 and 4), and evaluates every real `m.c13` constraint row at two adversarial check patterns.

Table 6 confirms both lemmas computationally, and shows the two formulations fail in *complementary, non-overlapping* cases, exactly as predicted in Section 4.7: when checks occur at both window endpoints, all 42 real split rows are satisfied even though 1 200 minutes are flown against the 600-

Table 6: Numeric verification of Lemmas 1 and 2 against the real `m.c13` constraints (both formulations) and the exact-state recursion (`results/tables/c13_loophole_validation.csv`). Both scenarios fly 1200 real minutes in window  $(1, 4]$  against  $T_{\max}^A = 600$ .

| Check pattern                      | Split C13              | Paper C13              | Exact-state |
|------------------------------------|------------------------|------------------------|-------------|
| Both endpoints checked (days 1, 4) | <b>accepts (wrong)</b> | rejects (correct)      | rejects     |
| Only end-day checked (day 4)       | rejects (correct)      | <b>accepts (wrong)</b> | rejects     |

minute limit (Lemma 2), while the single paper-form row correctly rejects it. When only the end day is checked, the paper form’s single row is satisfied despite the same 1200-minute violation (Lemma 1), while all 42 split rows correctly reject it (one row is violated). In both scenarios the exact-state recursion (17)–(21), evaluated on the same flight and check data, correctly flags the violation. This is a computational proof, not merely an algebraic one: the values above come from evaluating the actual Pyomo constraint objects built by `src/model.py`, with its real  $M_{\text{BIG}} = 9\,999\,999$  constant, rather than a hand re-derivation of the inequalities.

#### 7.5. Legacy hierarchy diagnostics and per-check-type effectiveness

Table 7 reports `experiments/hierarchy_effectiveness.py`: for each check type in isolation (A, B, C, D) we solve the legacy endpoint-split MILP on `ABCD_near_threshold_test`, and for the full mixed hierarchy on `ABCD_check_hierarchy_test`. Results are cross-validated with two independent exact solvers: the open-source HiGHS solver (`appsi_highs`) and the full, unrestricted IBM ILOG CPLEX 12.6 build distributed with AMPL (invoked through Pyomo’s `cplexamp` ASL interface); both report identical

objectives and statuses on every instance in this study. The CPLEX copy already on this machine’s PATH is a size-restricted Community/Preview edition (1 000 variable/constraint limit, **CPLEX Error 1016**) that cannot solve any instance with maintenance constraints active; it is not used for any reported result. Every solved instance is independently re-verified by replaying its solution: hour-triggered checks (A, B) are checked with the exact-state recursion of Section 4.8, and calendar-triggered checks (C, D) are checked with the analogous day-gap recursion, both implemented from scratch in the validator, not reused from the solver’s own constraints.

This table evaluates the legacy hierarchy implementation only. The event-based model is not included because this experiment does not solve and replay a matched event-based counterpart on the same hierarchy instances. Adding an event column without that matched evidence would create a misleading comparison; cross-formulation evidence is reported separately in Table 9.

For each check type solved in isolation, the legacy model returns an optimal, fully covered schedule that the independent replay certifies as correct: zero threshold violations across all four check types on `ABCD_near_threshold_test`. This is the practical evidence that the corrected hierarchy model enforces every check type (hour-triggered A/B and calendar-triggered C/D) without relaxation on realistic, non-adversarial instances.

On the full mixed-hierarchy instance, the replay flags four aircraft/check violations, of which three (aircraft 0’s A and B checks, aircraft 1’s B check) are fully explained by the instance’s deliberately overdue initial state ( $\Phi_0^A = 27\,001 > T_{\max}^A = 27\,000$ , similarly for  $B$ ) – a data precondition violation that would also invalidate the exact-state model’s own initial condition (22),

Table 7: Legacy-model per-check-type effectiveness on `ABCD_near_threshold_test` (single check active) and the full hierarchy on `ABCD_check_hierarchy_test` (`results/tables/hierarchy_effectiveness.csv`). “Genuine violations” excludes cases fully explained by an aircraft that starts the horizon already over threshold ( $\Phi_j > T_{\max}^c$ ), which violates the model’s own initial-condition precondition (22) regardless of formulation.

| Check(s)     | Status  | Obj    | Vars  | Cons. | Independently valid? | Genuine viol. |
|--------------|---------|--------|-------|-------|----------------------|---------------|
| A only       | optimal | 20 000 | 648   | 1 318 | yes                  | 0             |
| B only       | optimal | 20 000 | 648   | 1 516 | yes                  | 0             |
| C only       | optimal | 20 500 | 648   | 1 679 | yes                  | 0             |
| D only       | optimal | 20 500 | 648   | 1 679 | yes                  | 0             |
| Full A/B/C/D | optimal | 14 700 | 1 092 | 2 445 | <b>no</b>            | 1 of 4        |

independent of which C13 formulation is used. The remaining case is genuine: aircraft 2 starts with  $\Phi_2^A = 26\,800 < T_{\max}^A$  (not overdue), yet the legacy solution flies enough on day 1 to reach 27 200 minutes, exceeding the 27 000-minute threshold, a violation the legacy split model’s own constraints do not detect. This is a second, independently arising instance of the same endpoint-split flaw proved in Lemma 2 and Table 6 – found on an existing test instance under the full hierarchy, not constructed by hand – and it demonstrates that the flaw is not confined to contrived corner cases: it can occur whenever the hierarchy’s aggregated mega indicator lets a heavier check on one aircraft satisfy the endpoint condition for another aircraft’s window in the same solve.

### 7.6. Measured effectiveness of the event-based reformulation

Section 6 proves the event-based reformulation exact and immune to the endpoint-split failure mode by construction. This subsection measures its practical effectiveness against the day-indexed legacy model on which it improves, using `experiments/compare_formulations.py`. While diagnosing the first solved comparisons below, the same protocol used to certify results elsewhere in this section – always cross-checking a status against an independent computation – exposed an implementation defect in `src/event_model.py`: the pre-horizon initial-hours constraint used a big- $M$  equal to the source flight’s own duration plus its initial hours, which is not a valid upper bound on  $u_{ijk}$  once a route has accumulated several flights, and made the model reject cost-optimal schedules it should accept. The fix (using  $T_k + p_i$ , the same bound family as Equation (C12)) is included in the code accompanying this paper; all results below use the corrected version.

*Solved comparison on a well-posed instance..* On a small hand-built instance with two aircraft, six flights, and a threshold comfortably above the total flying time (`data/instances/event_vs_legacy_clean_test.json`), both formulations now return the same optimal objective (600, a single aircraft covering all six flights with no maintenance event), using CPLEX 12.6 for AMPL as the solver of record. This is exact agreement between two independently coded formulations on a genuinely solved instance, not merely a build-time comparison.

*A modelling-convention gap on the existing ABCD suite..* On the existing `ABCD_*` boundary instances (for example `ABCD_capacity_bottleneck_test`),

the event-based model reports `infeasible` where the legacy model reports `optimal`. This is not a solver or implementation defect: these instances’ own design notes (e.g. `ABCD_capacity_bottleneck_test.json`) state that an aircraft is intended to fly *past* its hour threshold and only be checked that same night – exactly the permissive convention encoded in the legacy model’s C13b relaxation. The exact-state formulations of this paper (both the day-indexed recursion of Section 4.8 and the event-based Equation (C11) bound  $u_{ijk} \leq T_k x_{ij}$  unconditionally) instead require the threshold to never be exceeded at any arrival, check or no check. Under that stricter, arguably safer reading, these particular instances are genuinely infeasible: no schedule exists in which an aircraft never exceeds its hour limit for these initial states and flight durations. This is a modelling-convention gap between the existing benchmark suite and the corrected formulations of this paper, not a defect in either implementation, and it is a direct practical consequence of fixing the endpoint-split flaw: the corrected model can no longer quietly accept the same after-the-fact resolution the legacy convention relied on.

*Model-size scaling versus planning horizon..* Table 8 reports `-build-only` statistics for both formulations on the same fleet (10 aircraft) across horizons  $h \in \{7, 15, 21, 30\}$  from the density-1.0 `DataCplex` grid.

The variable-count advantage predicted by Table 4 holds at  $h = 7, 15, 21$  (27–36% fewer variables) but reverses at  $h = 30$ , where the event model uses 20% *more* variables. Constraint count is larger for the event-based model at every horizon tested, by a factor of 1.6–2.8, and this factor broadly increases with  $h$ . The reason is exactly the caveat already noted in Section 6: constraint (C12) is built over every feasible connection  $(i, \ell) \in \mathcal{E}$ , and  $|\mathcal{E}|$  grows

Table 8: Model size versus planning horizon  $h$  (10 aircraft, density 1.0; `results/tables/formulation_scaling.csv`). “Vars ratio”/“Cons. ratio” are event-based over legacy.

| $h$ | Legacy vars | Event vars | Vars ratio | Legacy cons. | Event cons. | Cons. ratio |
|-----|-------------|------------|------------|--------------|-------------|-------------|
| 7   | 10 800      | 7 866      | 0.73       | 15 819       | 24 874      | 1.57        |
| 15  | 33 400      | 23 582     | 0.71       | 42 927       | 91 598      | 2.13        |
| 21  | 61 500      | 39 398     | 0.64       | 78 394       | 167 211     | 2.13        |
| 30  | 175 430     | 210 161    | 1.20       | 432 700      | 1 209 411   | 2.80        |

super-linearly with the number of flights on these dense ( $\delta = 1.0$ ) synthetic instances, whereas the legacy model’s maintenance constraint count scales with  $n_p n_d$ . The event-based reformulation is therefore not unconditionally smaller in practice: its variable-count advantage is real but bounded, and its constraint count is measurably larger on this benchmark family, a genuine trade-off rather than a strict improvement.

An exploratory density-0.5 horizon stress run is retained in the companion repository as `results/tables/formulation_horizon_scaling.csv`, but is excluded from runtime claims because its infeasible and time-limited rows mix different feasible domains. The corrected independent replay also showed that the legacy endpoint-split optimum on the  $h=10$  exploratory instance can violate the A/B flight-hour state limits. The controlled runtime comparison below therefore uses only feasible heavy instances for which both formulations are optimal and objective-matched.

The matched stress comparison also uncovered and corrected an objective-model defect in the legacy implementation. The day-indexed model repre-

sents a multi-day check by consecutive occupancy variables and previously charged a flat penalty per occupied day. When a check extended beyond the planning horizon, this underpriced the event relative to the event formulation, which charges once per check using its full duration. The legacy model now includes an explicit binary maintenance-start variable and charges the same duration-based event cost as the event formulation. After this correction, the  $h = 21$  heavy case has objective 881 263 in both formulations; the corrected matched result is recorded in `results/tables/formulation_h21_objective_aligned_v2.csv`.

*A guaranteed-feasible instance generator..* Both formulations are exact, so any solved-objective comparison between them is only informative on instances that are actually feasible; the random `DataCplex` generator can produce infeasible cells at higher densities, which is uninformative noise for this specific comparison rather than a property of either formulation. `src/generate_feasible_instances.py` instead builds each aircraft an explicit hub-and-spoke rotation (spoke  $\rightarrow$  maintenance airport  $\rightarrow$  spoke) with maintenance duration and turn-time slack built into every leg, and verifies every seeded rotation against an independent greedy timeline simulator before writing the instance, so the generated schedule is feasible by construction rather than by chance. Re-running the same  $h \in \{7, 15, 21, 30\}$ , 10-aircraft, density-1.0 comparison on four such instances, both formulations solve to a matching optimum on every instance. The generator caps flights per rotation, so the  $h \in \{15, 21, 30\}$  instances coincide in flight count (and hence in model size) and differ only by random seed; they test feasibility at equal size rather than a growing horizon:

This is the cleanest evidence in this paper for the correctness claim of

Table 9: Solved comparison on guaranteed-feasible instances (`results/tables/formulation_feasible_scaling.csv`), full CPLEX for AMPL, 180s time limit. Both formulations reach the identical objective on every instance. The  $h \in \{15, 21, 30\}$  rows share model size because the feasible generator caps flights per rotation; they differ only by random seed.

| Instance | Legacy status | Legacy obj. | Event status | Event obj. | Legacy vars/cons. | Event vars/cons. |
|----------|---------------|-------------|--------------|------------|-------------------|------------------|
| h=7      | optimal       | 24 935      | optimal      | 24 935     | 3 980/7 928       | 6 865/30 3       |
| h=15     | optimal       | 29 914      | optimal      | 29 914     | 7 020/15 494      | 8 820/38 7       |
| h=21     | optimal       | 29 886      | optimal      | 29 886     | 7 020/15 494      | 8 820/38 7       |
| h=30     | optimal       | 29 875      | optimal      | 29 875     | 7 020/15 494      | 8 820/38 7       |

Proposition 4: on four independently constructed, genuinely feasible instances, two independently coded exact formulations – built from different variable sets, different constraint families, and (after the fix above) different big- $M$  derivations – agree exactly on the optimal cost every time.

*Runtime and model-size scaling.* To complement the objective-parity tables, we measured the corrected legacy and event formulations on three independently generated feasible heavy instances with  $h \in \{7, 15, 21\}$ , ten aircraft, and density 0.5. Every run reported here was optimal, and the objective values matched exactly. The figures use wall-clock time, active Pyomo variable count, and active constraint count from `results/tables/formulation_heavy_scaling_optimize`. They therefore compare execution and model size only on a common, solved feasible domain.

The event formulation is therefore not claimed to be faster on this sample: its solve wall time is approximately 6.74, 1.68, and 1.43 times the legacy wall




Figure 1: Wall-clock solve time versus horizon on objective-matched feasible heavy instances. The event formulation is slower on this sample, although both formulations return the same optimum.

time at  $h = 7, 15, 21$ , respectively, in this benchmark run. A decrease does appear in model-construction time at the two longer horizons: the event/legacy build-time ratios are 1.13, 0.84, and 0.69. Thus the event formulation reduces build time after  $h = 7$  in this sample, but does not reduce total solved execution time. These are single-run measurements and should be interpreted as indicative rather than universal performance constants. Its advantage is structural correctness and a smaller variable set on these instances, while the additional connection constraints create a real computational trade-off. This distinction is important: the event formulation improves the invalid maintenance-state logic, but it is not a universal runtime replacement for the day-indexed model.




Figure 2: Active variable count versus horizon on the same feasible heavy instances. The event formulation uses fewer variables in these three cases.

*A second implementation defect, found and fixed the same way..* The same generator, run with `-maintenance-families ABCD`, initially solved under the legacy model but returned `infeasible` under the event-based model. Applying the same fault-isolation procedure used for Equation (C13) (relax all binary variables, solve the LP, and deactivate one named constraint block at a time) to constraints (C14a)–(C14b) found a second, distinct defect: the candidate set for a calendar-check deadline was built from *cost-feasible* assignment arcs only, without checking whether the candidate flight is actually *reachable* by that aircraft’s route (a flight can be affordable for an aircraft under the cost matrix’s ferry penalty while having no predecessor path to it at all in this sparse hub-and-spoke instance family). The deadline constraint could then require a check on a flight the aircraft could never legally reach, an




Figure 3: Active constraint count versus horizon on the same feasible heavy instances. The event formulation uses more constraints because its flow equations are generated over feasible event connections.

unsatisfiable requirement masquerading as a routing infeasibility. A second, narrower defect surfaced during the fix: the placeholder row inserted when no candidate exists at all evaluated to a plain Python `False` instead of a symbolic Pyomo constraint, raising a solver-side error rather than a clean infeasible status. Both are fixed in the code accompanying this paper (an explicit reachability check via the same predecessor sets used by constraint (C2), and a placeholder anchored on a real variable). Root-causing the resulting failure on the regenerated instance traced to a third defect, this time in the instance generator itself: the calendar-check deadline slack was computed from the coarse departure-day index rather than the outbound flight’s actual arrival time, so the day-rounding in `src/generate_feasible_instances.py` could place an aircraft’s own designated check-triggering flight a few minutes past




Figure 4: Relative event-based time on the objective-matched heavy instances. A ratio below one means the event formulation is faster; a ratio above one means it is slower. Build time can decrease at longer horizons, but solved wall time remains above the legacy value in every tested case.

its own deadline. With all three fixes applied, the regenerated ABCD instance solves **optimal** under both formulations with the identical objective (10 327), completing the cross-validation started above. We report the diagnostic process itself, not only the fix, because it is the same generic technique (LP relaxation plus single-block deactivation) used earlier for Constraint (13), applied here to a different constraint family with a different, unrelated root cause.

### 7.7. Correctness comparison of the maintenance formulations

The central advantage of the event-based formulation is correctness of the maintenance state, not feasibility on historical instances designed around a

Table 10: Correctness and objective-parity comparison of the maintenance formulations. The first two rows are the adversarial C13 scenarios from Table 6; “accepts” means the built model constraints admit the violating point, and “rejects” means they do not. The final rows are feasible heavy instances solved with unrestricted CPLEX 12.6 for AMPL and a 60-second limit.

| Test                        | Legacy paper C13 | Legacy split C13 | Event exact-state | Objective parity      |
|-----------------------------|------------------|------------------|-------------------|-----------------------|
| C13: both endpoints checked | rejects          | accepts (wrong)  | rejects           | correct rejection     |
| C13: only end checked       | accepts (wrong)  | rejects          | rejects           | correct rejection     |
| Feasible heavy, $h = 7$     | –                | optimal          | optimal           | 827 858 = 827 858     |
| Feasible heavy, $h = 15$    | –                | optimal          | optimal           | 1 132 028 = 1 132 028 |
| Feasible heavy, $h = 21$    | –                | optimal          | optimal           | 881 263 = 881 263     |

more permissive convention. Table 10 therefore compares the formulations on two adversarial Constraint (13) scenarios and on three independently generated feasible heavy instances. The first two rows evaluate the actual production constraints and the exact-state replay on the same 1 200-minute violation; the last three rows test whether the stricter event state changes the optimal assignment when a feasible schedule exists.

The table shows the intended result directly. The original paper C13 and the endpoint-split strengthening each fail on one complementary adversarial pattern, whereas the event formulation rejects both violations. On feasible heavy instances, this stronger state logic does not sacrifice the optimum: the event and corrected legacy implementations return identical objective values at all three horizons. Historical boundary cases that rely on after-threshold checking are excluded from this correctness table because they test a different maintenance convention rather than the quality of the event reformulation.

### 7.8. *Summary of findings*

1. The current runs validate the corrected formulations on deterministic synthetic instances; they do not reproduce the published data or averages.
2. The Constraint (13) endpoint-split flaw of Lemma 2 is confirmed computationally: all 42 real constraint rows built by `src/model.py` for the loophole instance are satisfied at an adversarial point where 1 200 minutes are actually flown against a 600-minute limit, while the exact-state recursion correctly rejects it.
3. Solved in isolation, each of the four check types (A, B, C, D) is enforced correctly by the legacy hierarchy model on `ABCD_near_threshold_test`: optimal, fully covered, and independently validated with zero violations.
4. Under the full mixed hierarchy, the same flaw reappears unprompted on an existing test instance (aircraft 2 in `ABCD_check_hierarchy_test` exceeds its A-check threshold by 200 minutes with no violated constraint reported), showing the issue is not confined to hand-built corner cases.
5. The correctness comparison shows that the event formulation rejects both adversarial C13 violations, while the original paper row and the endpoint-split legacy formulation each fail on one complementary pattern.
6. After correcting legacy multi-day maintenance-event pricing, the legacy and event formulations reach identical objectives on all three feasible heavy instances reported in Table 10.

7. Historical boundary cases that depend on after-threshold checking are excluded from the correctness comparison because they use a different maintenance convention; dense exploratory cases remain solver diagnostics.
8. The CPLEX installation available in this environment is a size-restricted Community/Preview edition and cannot solve any instance with maintenance constraints active (CPLEX Error 1016); the new Constraint (13) and hierarchy results therefore use the open-source HiGHS solver, which is exact and unrestricted.

## 8. Conclusions and Future Work

This paper revisited the compact Tail Assignment Problem model of Khaled et al. (2018), making four concrete contributions.

*C2–C4 feasibility correction..* We showed that the basic balance constraints can admit a simultaneous- departure corner case that is mathematically feasible but operationally impossible. Adding explicit pairwise non-overlap constraints restores physical executability without altering the variable set of the compact model. This correction also clarifies the domain on which maintenance constraints are evaluated.

*Constraint (13) correction..* We proved that the original cumulative-flight-time constraint can leave the opening maintenance interval unbounded, and that the natural two-row endpoint split is also insufficient in the complementary case where both window endpoints carry a check. Both lemmas are confirmed numerically against the real built constraint objects (Section 7.4):

each formulation accepts a schedule the other correctly rejects, on the same 1 200-minute violation against a 600-minute limit. The exact replacement uses a continuous cumulative state with maintenance-triggered resets and explicit initial conditions, given first in day-indexed form (Section 4.8) and then, completely independently of any day index, in the event-based reformulation of Section 6.

*Effect on C8–C14.* The non-overlap correction changes the assignment domain on which C8–C14 operate but does not repair cumulative-hour propagation by itself. Theorem 1 therefore relies specifically on the exact state constraints, together with the corrected basic feasibility and maintenance-event constraints.

*Full  $\{A, B, C, D\}$  hierarchy.* We extended the single-type framework to an explicit four-level A/B/C/D hierarchy. Effective-check indicators encode heavier-check resets, separate start and occupancy variables represent multi-day C/D checks, and initial hour/day values carry maintenance history into the horizon. Each check type solved in isolation, and the full hierarchy together, is independently re-verified against the raw flight timetable (Section 7.5); the full-hierarchy run also reproduces the Constraint (13) split flaw unprompted on an existing test instance, showing the issue is not confined to hand-built corner cases.

*Event-based reformulation.* We gave a complete, self-contained, day-free reformulation of the same C1–C14 requirements (Section 6), in which maintenance is a virtual flight indexed on its triggering flight rather than on calendar day, and proved its correctness directly (Proposition 4). Because its flight-

hour counter propagates arc by arc along the aircraft’s actual route rather than over a day-pair window, it cannot exhibit the split-form failure mode of Constraint (13) by construction. Measuring this reformulation against the day-indexed model (Section 7.6) surfaced and let us fix three genuine implementation defects using the same generic technique (LP relaxation plus single-constraint-block deactivation): a big- $M$  sizing defect in the flight-hour counter, a missing routing-reachability check in the calendar-check deadline constraint, and an insufficient day-rounding margin in the instance generator itself, after which both formulations agree exactly on every guaranteed-feasible test instance, including one exercising the full  $\{A, B, C, D\}$  hierarchy. Measurement also revealed that the event-based model’s variable-count advantage is real but not universal (it reverses at the largest horizon tested) and that its constraint count is consistently larger on dense timetables, and that its stricter, never-exceed-the-threshold semantics is a genuine modelling-convention change from the existing benchmark suite, not merely an implementation detail. We report these as honest, measured trade-offs rather than an unconditional improvement.

*Reproducibility..* The project contains the code, deterministic instance generator, experiment configurations, and result tables needed to rerun the reported snapshots. The legacy split formulation and the separate event/state formulation are retained as distinct implementations so their future performance can be compared without conflating their maintenance semantics.

*Threats to validity..* The reported computational comparisons are intentionally scoped to the synthetic benchmark family used in this study. The quick grids use few instances, and several dense runs end in infeasible/error states.

The paper-versus-split-versus-event-based C13 comparison is completed for a targeted counterexample and for the per-check-type hierarchy study, but not yet at full paper-scale grid sizes. Performance conclusions are therefore descriptive of the reported runs only. Decomposition-heavy and robust/scenario alternatives may exhibit different trade-offs under their native assumptions and tuning protocols.

#### *Future directions*

- **Integrated fleet assignment and routing:** the compact model is well-suited for integration with fleet-assignment decisions, as noted by Khaled et al. (2018). The absence of the cyclic constraint facilitates this integration.
- **Robust and recoverable variants:** extending Borndörfer et al. (2010); Froyland et al. (2013) to the compact formulation with full maintenance hierarchy is an open problem.
- **Controlled formulation benchmark at scale:** the paper-C13, endpoint-split, and event-based builders are now separately selectable and every solution is independently validated by replaying its maintenance counters; extending this comparison of model size, relaxation bounds, runtime, and feasibility to the full paper-scale instance grid remains open.
- **Heuristic and metaheuristic methods:** constructive and learning-based heuristics for the corrected, full-hierarchy model are deliberately out of scope here and are left for a companion study.

## References

- Barnhart, C., Belobaba, P., Odoni, A.R., 2003. Applications of operations research in the air transport industry. *Transportation Science* 37, 368–391.
- Barnhart, C., Boland, N.L., Clarke, L.W., Johnson, E.L., Nemhauser, G.L., Shenoi, R.G., 1998. Flight string models for aircraft fleetings and routing. *Transportation Science* 32, 208–220. doi:10.1287/trsc.32.3.208.
- Borndörfer, R., Dovica, I., Nowak, I., Schickinger, T., 2010. Robust tail assignment. Technical Report. Konrad-Zuse-Zentrum für Informationstechnik Berlin.
- Clarke, L., Johnson, E., Nemhauser, G., Zhu, Z., 1997. The aircraft rotation problem. *Annals of Operations Research* 69, 33–46. doi:10.1023/A:1018945415148.
- Clarke, L.W., Hane, C.A., Johnson, E.L., Nemhauser, G.L., 1996. Maintenance and crew considerations in fleet assignment. *Transportation Science* 30, 249–260.
- Cordeau, J.F., Stojković, G., Soumis, F., Desrosiers, J., 2001. Benders decomposition for simultaneous aircraft routing and crew scheduling. *Transportation Science* 35, 375–388.
- Desaulniers, G., Desrosiers, J., Dumas, Y., Solomon, M.M., Soumis, F., 1997. Daily aircraft routing and scheduling. *Management Science* 43, 841–855.
- Froyland, G., Maher, S.J., Wu, C.L., 2013. The recoverable robust tail assignment problem. *Transportation Science* 48, 351–372.

- Gopalan, R., Talluri, K.T., 1998. Mathematical models in airline schedule planning: A survey. *Annals of Operations Research* 76, 155–185.
- Grönkvist, M., 2005. The tail assignment problem. Ph.D. thesis. Chalmers, Göteborg University.
- Hane, C.A., Barnhart, C., Johnson, E.L., Marsten, R.E., Nemhauser, G.L., Sigismondi, G., 1995. The fleet assignment problem: Solving a large-scale integer program. *Mathematical Programming* 70, 211–232.
- Haouari, M., Shao, S., Sherali, H.D., 2012. A lifted compact formulation for the daily aircraft maintenance routing problem. *Transportation Science* 47, 508–525.
- Hart, W.E., Laird, C.D., Watson, J.P., Woodruff, D.L., Hackebeil, G.A., Nicholson, B.L., Sirola, J.D., 2023. Pyomo – Python-based Optimization Modeling Objects. Version 6.x, <https://www.pyomo.org/>.
- Huangfu, Q., Hall, J.A.J., 2018. Parallelizing the dual revised simplex method. *Mathematical Programming Computation* 10, 119–142. HiGHS solver, <https://highs.dev/>.
- IBM, 2023. IBM ILOG CPLEX Optimization Studio: CPLEX User’s Manual. Version 22.1.
- Khaled, O., Minoux, M., Mousseau, V., Michel, S., Ceugniet, X., 2018. A compact optimization model for the tail assignment problem. *European Journal of Operational Research* 264, 548–557. doi:10.1016/j.ejor.2017.06.045.

- Klabjan, D., 2005. Large-scale models in the airline industry, in: Desaulniers, G., Desrosiers, J., Solomon, M.M. (Eds.), *Column Generation*. Springer US, pp. 163–195.
- Lacasse-Guay, E., Desaulniers, G., Soumis, F., 2010. Aircraft routing under different business processes. *Journal of Air Transport Management* 16, 258–263.
- Lavoie, S., Minoux, M., Odier, E., 1988. A new approach for crew pairing problems by column generation with an application to air transportation. *European Journal of Operational Research* 35, 45–58.
- Levin, A., 1971. Scheduling and fleet routing models for transportation systems. *Transportation Science* 5, 232–255.
- Liang, Z., Chaovalitwongse, W.A., 2012. A network-based model for the integrated weekly aircraft maintenance routing and fleet assignment problem. *Transportation Science* 47, 493–507.
- Liang, Z., Chaovalitwongse, W.A., Huang, H.C., Johnson, E.L., 2011. On a new rotation tour network model for aircraft maintenance routing problem. *Transportation Science* 45, 109–120.
- Mercier, A., Cordeau, J.F., Soumis, F., 2005. A computational study of Benders decomposition for the integrated aircraft routing and crew scheduling problem. *Computers & Operations Research* 32, 1451–1476. doi:10.1016/j.cor.2003.11.013.

- Minoux, M., 1984. Column generation techniques in combinatorial optimization, a new application to crew pairing problems, in: Proceedings of the AGIFORS Symposium, Strasbourg, France.
- Sarac, A., Batta, R., Rump, C.M., 2006. A branch-and-price approach for operational aircraft maintenance routing. *European Journal of Operational Research* 175, 1850–1869.
- Sherali, H.D., Bish, E.K., Zhu, X., 2006. Airline fleet assignment concepts, models, and algorithms. *European Journal of Operational Research* 172, 1–30.