

# Integrated Model for Aircraft Routing and Maintenance Planning

Thesis Manuscript

August 18, 2026

# Abstract

This thesis studies the integrated aircraft routing and maintenance planning problem within a compact mixed-integer programming framework. We begin from the compact tail-assignment model and show that the basic feasibility conditions are insufficient for physical realism because they do not fully exclude simultaneous departures and overlapping aircraft trajectories. We then formalize and correct the cumulative flight-hour maintenance constraint, which is the central modeling issue in the original formulation. In addition, the thesis extends the model to a full four-level maintenance hierarchy using A/B/C/D checks, while maintaining the structure of a compact and operationally meaningful assignment model. The resulting formulation is evaluated against constructive and search-based heuristics and benchmarked on generated instance families designed to capture varying density and maintenance pressure. The study shows that the corrected maintenance logic yields valid schedules, tighter formulations, and a more reliable comparison between exact and heuristic strategies.

# Acknowledgements

The author acknowledges the value of the underlying modeling and computational work that motivated this thesis. The research is built on the existing compact tail-assignment formulation and its maintenance extension, with emphasis on mathematical correctness, operational feasibility, and reproducible computational evidence.

# Contents

|          |                                                                         |           |
|----------|-------------------------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                                     | <b>1</b>  |
| 1.1      | Introduction . . . . .                                                  | 1         |
| <b>2</b> | <b>Literature Review</b>                                                | <b>4</b>  |
| 2.1      | Literature Review . . . . .                                             | 4         |
| 2.1.1    | Aircraft routing and tail-assignment formulations . . . . .             | 4         |
| 2.1.2    | Mathematical programming formulations . . . . .                         | 5         |
| 2.1.3    | MILP alternatives to the Khaled compact model . . . . .                 | 6         |
| 2.1.4    | Integrated airline planning . . . . .                                   | 7         |
| 2.1.5    | Robustness and operational uncertainty . . . . .                        | 7         |
| 2.1.6    | Maintenance planning at different decision levels . . . . .             | 7         |
| 2.1.7    | Heuristic solution methods and cross-domain transfer . . . . .          | 8         |
| 2.1.8    | Comparison scope and fairness . . . . .                                 | 9         |
| 2.1.9    | Research gap and position of the thesis . . . . .                       | 9         |
| <b>3</b> | <b>The Ground Routing Model</b>                                         | <b>10</b> |
| 3.1      | Scope and modelling assumptions . . . . .                               | 10        |
| 3.2      | Sets, parameters, and cumulative flight sets . . . . .                  | 11        |
| 3.3      | Compact mixed-integer formulation . . . . .                             | 11        |
| 3.3.1    | C1: flight coverage . . . . .                                           | 12        |
| 3.3.2    | C2: continuity at non-initial airports . . . . .                        | 12        |
| 3.3.3    | C3: continuity at the initial airport . . . . .                         | 13        |
| 3.4      | Interpretation and validity of the balances . . . . .                   | 13        |
| 3.5      | Equal-time departure counterexample . . . . .                           | 13        |
| 3.6      | Conflict graph and corrected feasibility model . . . . .                | 14        |
| 3.7      | Model size and computational implications . . . . .                     | 15        |
| 3.8      | Illustrative non-cyclic assignment . . . . .                            | 15        |
| 3.9      | Chapter summary . . . . .                                               | 16        |
| <b>4</b> | <b>Maintenance-Extended Model and the Correction of Constraint (13)</b> | <b>17</b> |
| 4.1      | Scope and operational assumptions . . . . .                             | 17        |

|          |                                                         |           |
|----------|---------------------------------------------------------|-----------|
| 4.2      | Sets, parameters, and auxiliary definitions . . . . .   | 18        |
| 4.3      | Maintenance decision variables . . . . .                | 18        |
| 4.4      | Extended objective function . . . . .                   | 19        |
| 4.5      | Maintenance-event feasibility . . . . .                 | 20        |
| 4.5.1    | C8: blocking subsequent same-day flights . . . . .      | 20        |
| 4.5.2    | C9: coupling a check to its trigger flight . . . . .    | 20        |
| 4.5.3    | C10: airport-day maintenance capacity . . . . .         | 20        |
| 4.5.4    | C11 and C14: daily aggregation and uniqueness . . . . . | 20        |
| 4.5.5    | C12: maximum calendar spacing . . . . .                 | 21        |
| 4.6      | Original cumulative-hour constraint . . . . .           | 21        |
| 4.7      | Corrected endpoint-anchored formulation . . . . .       | 22        |
| 4.8      | Pre-horizon accumulated hours . . . . .                 | 23        |
| 4.9      | Choice of the big- $M$ coefficient . . . . .            | 23        |
| 4.10     | Physical interpretation and model validity . . . . .    | 24        |
| 4.11     | Complete single-type formulation and size . . . . .     | 24        |
| <b>5</b> | <b>Full A/B/C/D Maintenance Hierarchy</b>               | <b>26</b> |
| 5.1      | Check types and regulatory hierarchy . . . . .          | 26        |
| 5.2      | Hierarchy constraints . . . . .                         | 27        |
| 5.3      | Flight-hour checks: types A and B . . . . .             | 27        |
| 5.4      | Calendar-day checks: types C and D . . . . .            | 27        |
| 5.5      | Pre-horizon state accounting . . . . .                  | 28        |
| 5.6      | Complexity and model implications . . . . .             | 28        |
| <b>6</b> | <b>Heuristic Solution Methods</b>                       | <b>29</b> |
| 6.1      | Timeline-based feasibility oracle . . . . .             | 29        |
| 6.2      | Greedy construction . . . . .                           | 29        |
| 6.3      | Insertion-based repair . . . . .                        | 30        |
| 6.4      | Bounded local search . . . . .                          | 30        |
| 6.5      | Space-time network search methods . . . . .             | 31        |
| 6.6      | Quality assessment framework . . . . .                  | 31        |
| <b>7</b> | <b>Computational Study</b>                              | <b>32</b> |
| 7.1      | Instance families and benchmark design . . . . .        | 32        |
| 7.2      | Effect of the formulation correction . . . . .          | 33        |
| 7.3      | Exact solver performance . . . . .                      | 33        |
| 7.4      | Heuristic performance and optimality gaps . . . . .     | 33        |
| 7.5      | LP lower-bound analysis . . . . .                       | 34        |
| 7.6      | Summary of findings . . . . .                           | 34        |



# List of Figures

# List of Tables

|     |                                                                                                                                       |    |
|-----|---------------------------------------------------------------------------------------------------------------------------------------|----|
| 3.1 | Notation for the compact ground-routing model. . . . .                                                                                | 12 |
| 3.2 | Six-flight timetable for the illustrative ground assignment. . . . .                                                                  | 16 |
| 4.1 | Additional notation for the type-A maintenance model. . . . .                                                                         | 19 |
| 4.2 | Activation of the original and corrected C13 rows when $Y_{jdd'}^o = 0$ . “Active”<br>means that $H_{jdd'} \leq T_{\max}^A$ . . . . . | 22 |



# Chapter 1

## Introduction

### 1.1 Introduction

The airline planning process is traditionally decomposed into a sequence of sub-problems: timetable design, fleet assignment, crew rostering, tail assignment, and recovery planning [2, 16]. The *tail assignment problem* (TAP)—assigning a specific aircraft (identified by its tail number) to each flight leg—is the penultimate planning step and is tightly coupled with aircraft maintenance scheduling. Poor tail assignment decisions directly inflate operational costs through unnecessary ferry (repositioning) flights, avoidable leasing decisions, suboptimal maintenance routing, and service disruptions.

Khaled et al. [20] proposed a compact 0–1 linear programming formulation for the TAP, featuring only  $O(n_f n_p)$  binary variables (where  $n_f$  is the number of flights and  $n_p$  the fleet size). This is a significant reduction over the classic multi-commodity flow model of Levin [24], which requires  $O(n_f^2 n_p)$  variables. Computational experiments reported in that paper show that instances with up to 1 500 flight legs and 40 aircraft can be solved to certified optimality within one hour on a standard workstation.

However, most practical studies still rely on cyclic or short-horizon settings (daily/weekly schedules) and often separate routing from maintenance decisions. For modern airline operations, this is restrictive: schedules are frequently adjusted for seasonality and event-driven demand, while aircraft must satisfy multi-level maintenance obligations across the same horizon. Treating routing and maintenance independently shifts feasibility issues into the operational window, where corrective actions are expensive and disruptive.

This thesis addresses that gap by revisiting and extending the compact TAP framework toward an integrated, season-level routing-and-maintenance planning approach. The objective is to produce physically executable aircraft rotations that satisfy maintenance requirements and improve planning certainty before schedule execution.

## Thesis contributions

This thesis makes four contributions that verify, correct, and extend the state of the art:

1. **Basic feasibility correction (C2–C4 / constraints (3)–(6)).** We show that the strict-time balance definition used in the compact basic model admits a simultaneous-departure corner case: one aircraft can be assigned to two flights departing at the same airport and same time without violating the balance inequalities. We formalize this counter-example and enforce physical executability by adding explicit pairwise non-overlap constraints.
2. **Constraint (13) correction.** We show that the cumulative-flight-time maintenance constraint (Constraint 13 in Khaled et al. [20]) contains a logical flaw: the big- $M$  linearisation term  $M(2 - y_{jd} - y_{jd'})$  leaves the constraint vacuously satisfied whenever at least one of the endpoint days  $d, d'$  hosts no check, so an accumulation window that is not bounded by checks at both ends is left unconstrained and may exceed the flight-hour limit. We provide a corrected formulation using two split endpoint constraints, prove its validity, and show that the corrected formulation yields strictly tighter LP relaxations on high-density instances.
3. **Full  $\{A, B, C, D\}$  maintenance hierarchy.** The model of Khaled et al. [20] handles type-A checks (cumulative flight hours). We extend the formulation to the full four-level hierarchy mandated by aviation regulators (FAA/EASA): type-A and B checks are flight-hour triggered; type-C and D checks are calendar-day triggered. A heavier check subsumes all lighter checks (hierarchical counter resets). We further add a constraint accounting for pre-horizon accumulated maintenance hours/days per aircraft and integrate these requirements into long-horizon schedule planning.
4. **Heuristic benchmark.** We implement and evaluate two constructive heuristics: (a) a greedy/insertion heuristic operating directly on flight lists, and (b) a Dijkstra-based heuristic on the space-time network representation (adapted from Grönkvist [17]). Both heuristics respect the full maintenance hierarchy. We report optimality gaps against the MILP on the same instance sets and show the Dijkstra heuristic consistently outperforms an ACO alternative.

## Thesis organisation

Chapter 2 reviews related work. Chapter 3 restates the basic TAP model and proves a C2–C4 corner-case correction. Chapter 4 presents the maintenance-extended model, the Constraint (13) correction, and the induced interpretation updates for C8–C14 under the corrected basic model. Chapter 5 describes the full  $\{A, B, C, D\}$  extension.

Chapter 6 analyses the heuristic algorithms. Chapter 7 reports all computational results. Chapter 8 concludes.

# Chapter 2

## Literature Review

### 2.1 Literature Review

The problems studied in the domain of thesis can be divided into the following classes:

- Fleet assignment - selection of an aircraft type for each flight;
- Aircraft routing - construction of maintenance-feasible sequences for a homogeneous fleet;
- Tail assignment - allocation of individual aircraft, with their own locations and maintenance states, to those sequences; and
- Maintenance planning - scheduling checks and the tasks performed during them.

This distinction matters because models from these streams cannot be compared solely by instance size or solution time. They differ in decision scope, maintenance semantics, planning horizon, and treatment of uncertainty [16, 2, 22].

In the following sections, we review the literature on these four problem classes, as well as on integrated airline planning, robustness and operational uncertainty, and heuristic solution methods. The review concludes with a discussion of the research gap addressed by this thesis.

#### 2.1.1 Aircraft routing and tail-assignment formulations

Early network formulations represent aircraft movement as flow through airport and time events. The multi-commodity model of Levin [24] tracks each aircraft separately, while the time-space network of Hane et al. [18] provides a natural representation of flight, ground, and overnight arcs. Maintenance opportunities can then be embedded in feasible rotations [7, 6]. These models make physical continuity explicit, but aircraft-specific connection variables grow rapidly with the number of feasible flight pairs.

Route-based formulations instead associate variables with complete feasible flight strings. The flight-string model of Barnhart et al. [1] exploits this structure through column generation. Such formulations can express detailed route feasibility in the pricing problem, but their implementation is substantially more involved than a directly solved compact MILP. Decomposition methods offer a related trade-off: they separate coupled decisions to obtain tractable subproblems, as illustrated by Benders approaches to aircraft routing and crew scheduling [8, 28].

Compact formulations avoid explicit route enumeration. The lifted daily maintenance-routing model of Haouari et al. [19], for example, strengthens a compact nonlinear formulation through reformulation and linearisation. Khaled et al. [20] pursue a different compact representation in which binary variables assign flights directly to individual aircraft and aggregate time-balance constraints impose route continuity without connection-arc variables. Its small variable count and direct solvability make it the closest baseline for this thesis. Compactness, however, transfers more logical work to aggregate constraints; consequently, corner cases and conditional maintenance bounds require particularly careful verification.

The daily model of Maher et al. [27] is an important operational counterpoint. It assigns aircraft to one-day lines of flight, enforces current maintenance needs, and uses look-ahead constraints for the following two days. Exact and heuristic procedures can modify lines of flight to reduce maintenance misalignment under day-to-day perturbations. That work addresses operational flexibility around a supplied route plan, whereas the present thesis audits and extends a deterministic compact assignment formulation over its full planning horizon. The two approaches are therefore complementary rather than direct competitors.

## 2.1.2 Mathematical programming formulations

Three main modelling paradigms underlie TAP research:

**Set partitioning / column generation.** The non-compact formulation enumerates *strings* (feasible flight sequences per aircraft) as 0–1 variables. Because the number of strings is exponential, column generation is required [23, 29, 1, 31]. Combining column generation with branch-and-bound (*branch-and-price*) gives exact solutions but at significant algorithmic overhead. Robust extensions appear in Borndörfer et al. [5] and Froyland et al. [15].

**Multi-commodity flow.** Levin [24] introduced the compact multi-commodity flow formulation with  $O(n_f^2 n_p)$  binary variables. This remains the standard benchmark for compact models and is the primary comparison point in Khaled et al. [20].

**Time-space networks.** Hane et al. [18] proposed the time-space network for the *fleet assignment problem*. Nodes represent airport–time events; arcs encode ground waits, flights, and overnight stays. Extensions to aircraft routing appear in Clarke et al. [7]. Preprocessing by node aggregation and island isolation reduces model size [18, 33]. Daily and weekly maintenance routing variants are studied in Liang and Chaovalitwongse [25] and Liang and Chaovalitwongse [26]; the latter also integrates fleet assignment decisions, but only approximate solutions are obtained via variable fixing.

**Nonlinear / lifted formulations.** Haouari et al. [19] present a compact formulation with nonlinear constraints that are linearised via the Reformulation–Linearisation Technique (RLT). The largest instances solved involve 344 flights, an order of magnitude below the instances handled by the model of Khaled et al. [20].

### 2.1.3 MILP alternatives to the Khaled compact model

Beyond the compact binary assignment model of Khaled et al. [20], several mixed-integer alternatives are relevant for exact TAP modelling and large-scale implementation:

- **Arc-flow MILP (time-indexed or event-indexed):** decision variables are attached to feasible flight-connection arcs instead of flight–aircraft assignment pairs, often yielding stronger network-structure constraints [24, 18, 33].
- **Route-based master MILP (set partitioning):** aircraft duties are represented as columns in a master MILP, typically solved with branch-and-price; this is exact but algorithmically heavier [1, 31, 21].
- **Decomposed MILP frameworks:** Benders- or Dantzig–Wolfe-style decompositions separate assignment/routing from maintenance-resource feasibility, improving scalability on larger fleets [8, 28, 11].
- **Two-stage or rolling-horizon MILP:** first assign flights, then repair/optimize maintenance and turnaround feasibility in a second MILP, trading global optimality for tractable runtime [7, 25, 26].
- **Robust/scenario MILP variants:** delay and disruption uncertainty is modelled through multiple scenarios with non-anticipativity-style linking constraints, producing schedules that are less brittle operationally [5, 15].

In this taxonomy, the Khaled model is best viewed as a compact, directly solvable baseline that is simple to implement and benchmark, while the alternatives above trade model size, decomposition complexity, and robustness for stronger scalability or richer operational realism.

### 2.1.4 Integrated airline planning

Sequential airline planning can propagate poor or infeasible decisions from fleet assignment to routing and then to crew planning. Shao et al. [32] integrate fleet assignment, aircraft routing, and crew pairing in a single framework and solve it using Benders decomposition with acceleration strategies. Their results demonstrate the value of preserving cross-stage dependencies, but also show why decomposition is needed once fleet, route, and crew decisions are combined.

At a narrower integration level, Mirjafari et al. [30] jointly model maintenance routing and crew scheduling, including aircraft replacement, deadheading, night-flight allocation, and flight-hour maintenance thresholds. They compare Lagrangian relaxation and particle swarm optimisation, with the relaxation producing solutions closer to the small-instance optimum. These integrated models establish that routing decisions interact with downstream resources. They do not, however, provide a like-for-like test of the compact tail-assignment constraints studied here, because their variables, objectives, and resource constraints cover a broader planning problem.

### 2.1.5 Robustness and operational uncertainty

Deterministic cost-optimal routes can be fragile when delays propagate through tight aircraft connections. Recoverable-robust tail assignment explicitly balances planned performance with the ability to repair a disrupted schedule [15]. A simulation-optimisation line of work instead uses stochastic evaluation to place connection buffers and retime flights. The initial particle-swarm and Monte Carlo framework of Ben Ahmed et al. [3] was developed into the weekly hybrid model of Ben Ahmed et al. [4]. The latter combines a linearised maintenance-routing model with simulation-based retiming and reports substantial improvements in on-time performance and propagated delay on airline data.

These studies reveal a limitation of the present thesis: the corrected model establishes deterministic physical and maintenance feasibility, but it does not by itself guarantee resilience to stochastic delays. Robust retiming, recourse, and scenario-based evaluation are therefore complementary extensions, not features that can be inferred from a feasible deterministic schedule.

### 2.1.6 Maintenance planning at different decision levels

Aircraft maintenance enters the literature at three levels that should not be conflated. First, routing models ensure that an aircraft reaches an eligible station before a usage limit is exceeded. Second, long-term check scheduling selects dates and capacities for maintenance packages. Third, task allocation assigns individual inspection and repair

tasks to those checks. The present thesis operates mainly at the first level, while using a stylised hierarchy to represent interactions among check categories.

At the long-term level, Deng et al. [9] use dynamic programming to coordinate A- and C-check schedules for a heterogeneous fleet while accounting for aircraft status and maintenance capacity. Their airline case study reports fewer checks and higher aircraft availability than current practice. At the task level, Witteman et al. [34] formulate repeated maintenance-task allocation as a time-constrained variable-sized bin-packing problem. Their worst-fit-decreasing heuristic obtains small gaps to an exact method while responding quickly enough for repeated replanning. Deng et al. [10] then integrate check scheduling, task allocation, and shift planning in a decision support system and demonstrate improvements on airline cases.

This evidence also requires more precise terminology than a simple statement that four checks are universally mandated. Letter checks are practical packages whose content and intervals depend on the approved maintenance programme, aircraft type, utilisation, and operator. The supplied maintenance studies note that B-check tasks are often absorbed into successive A-checks and that some operators combine D-check work with heavy C-checks [9, 34, 10]. Accordingly, the A/B/C/D structure in Chapter 5 is a modelling abstraction for nested usage limits and counter resets. It is not a claim that every operator implements four separate maintenance events.

For this thesis, the maintenance-planning literature has two further implications. Check feasibility depends on both utilisation counters, such as flight hours and cycles, and calendar limits; and an aircraft’s state before the planning horizon cannot be discarded. Conversely, a routing-level model that places a check does not solve hangar capacity, task packaging, workforce, or shift planning. Those resource layers remain outside the present formulation.

### **2.1.7 Heuristic solution methods and cross-domain transfer**

Constructive and local-search methods are operationally useful when exact solution times are incompatible with replanning. Grönkvist [17] develops tail-assignment heuristics in this setting. The original work targets a train application rather than flights, but the principles are transferable. The algorithm: represent movement and waiting in an acyclic space-time network, select locomotives in a cost-based order, repeatedly solve a least-cost path problem, update arc costs after assignment, and terminate a path when a maintenance limit becomes due. The algorithmic principles are applicable to aircraft because time also gives the routing network a natural topological order.

This transfer must preserve aircraft-specific feasibility. Airport continuity, turnaround time, aircraft eligibility, and all maintenance counters have to be encoded in nodes, arcs, or state labels before a shortest-path solution is a valid tail route. Dijkstra’s



shortest-path method provides the basic path engine [12]; ant colony optimisation provides a population-based alternative for exploring assignment sequences [14, 13]. In this thesis these methods are benchmarked as constructive alternatives to the compact MILP, not treated as evidence of optimality.

### 2.1.8 Comparison scope and fairness

For an empirical comparison to be methodologically fair, competing models should be evaluated under a common instance family, objective definition, maintenance semantics, and solver-budget protocol. In TAP, this condition is often not met across compact MILP, branch-and-price, Benders-decomposed, and robust/scenario formulations because these approaches typically optimize different decision scopes or uncertainty assumptions. Accordingly, this thesis uses Khaled et al. [20] as the direct compact baseline for instance-matched comparisons, and treats other MILP families as complementary reference paradigms rather than strict one-to-one benchmarks.

### 2.1.9 Research gap and position of the thesis

The reviewed literature offers sophisticated solutions for route generation, integrated airline planning, robust retiming, long-term check scheduling, and maintenance-task allocation. Less attention is paid to auditing the logical correctness of a compact published formulation before extending or benchmarking it. This thesis addresses that narrower but foundational gap using Khaled et al. [20] as the direct baseline.

The contribution is positioned in four parts. First, the aggregate ground-flow conditions are tested against simultaneous and overlapping flight patterns, and explicit non-overlap conditions are introduced where needed. Second, the conditional cumulative-hour logic of Constraint (13) is analysed and replaced by two endpoint-anchored inequalities. Third, the maintenance state is extended to multiple nested check categories, including pre-horizon usage, while clearly remaining at routing level rather than modelling maintenance tasks or workforce. Fourth, exact and heuristic methods are evaluated on common generated instances under the same objective and feasibility checks.

This scope supports a controlled claim: the thesis improves the correctness and maintenance coverage of one compact TAP formulation and evaluates practical solution alternatives for that formulation. It does not claim to supersede route-based, decomposition, robust, or maintenance-resource models, whose broader decision scopes answer different operational questions.

# Chapter 3

## The Ground Routing Model

This chapter develops the ground-routing component of the thesis. The model is based on the compact tail-assignment formulation of Khaled et al. [20], but maintenance decisions are deliberately excluded until Chapter 4. This separation isolates the physical conditions that make a flight assignment executable and ensures that the maintenance extension is built on valid aircraft trajectories.

The term *compact* refers to the variable space. Route-based models use variables for complete rotations, and arc-flow models use variables for feasible flight connections. The formulation below instead uses one binary variable for each flight–aircraft pair. It therefore has  $O(|\mathcal{F}||\mathcal{P}|)$  binary variables and can be submitted directly to a MILP solver. Airport continuity and timing are recovered through cumulative balance inequalities. This chapter derives those inequalities, identifies their equal-time departure limitation, and presents the corrected ground model used throughout the thesis.

### 3.1 Scope and modelling assumptions

The model is defined over a finite planning horizon using absolute times from the start of that horizon. A flight on day two consequently has larger time values than flights on day one. The following assumptions delimit the ground model:

1. every published flight leg must be covered exactly once;
2. each aircraft has one known initial airport;
3. an aircraft moves between airports only through an assigned flight, so any permitted ferry movement must be included explicitly;
4. departure and arrival times are fixed, and consecutive flights require at least the minimum ground turn time;
5. forbidden flight–aircraft pairs are removed or fixed to zero; and

6. maintenance and hangar decisions are deferred to Chapters 4 and 5.

The model is non-cyclic: an aircraft need not return to its initial airport by the end of the horizon. Its terminal position may be passed to the next planning period. Terminal-balance constraints can impose cyclicity, but this is an additional planning policy and may make a finite-horizon instance infeasible.

## 3.2 Sets, parameters, and cumulative flight sets

Let  $\mathcal{F} = \{1, \dots, n_f\}$  be the set of flight legs,  $\mathcal{P} = \{1, \dots, n_p\}$  the set of individual aircraft, and  $\mathcal{A} = \{1, \dots, n_k\}$  the set of airports. For flight  $i \in \mathcal{F}$ , let  $a_i^\uparrow$  and  $a_i^\downarrow$  be its departure and arrival airports, and let  $t_i^\uparrow$  and  $t_i^\downarrow$  be its departure and arrival times. The arrow notation follows the flight event:  $\uparrow$  denotes a departure and  $\downarrow$  an arrival. Parameter  $a_j^0$  is the initial airport of aircraft  $j$ ,  $\Delta t > 0$  is the minimum turn time, and  $c_{ij}$  is the cost of assigning flight  $i$  to aircraft  $j$ .

For airport  $k \in \mathcal{A}$  and threshold  $\theta$ , define

$$\mathcal{F}_k^\downarrow(\theta) = \{i \in \mathcal{F} : a_i^\downarrow = k, t_i^\downarrow \leq \theta - \Delta t\}, \quad (3.1)$$

$$\mathcal{F}_k^\uparrow(\theta) = \{i \in \mathcal{F} : a_i^\uparrow = k, t_i^\uparrow < \theta\}. \quad (3.2)$$

The arrival set contains flights that have reached  $k$  early enough to complete the turn before  $\theta$ . The departure set contains flights that have left  $k$  strictly before  $\theta$ . The strict inequality is central to the formulation and to the counterexample in Section 3.5.

The cumulative inventory of aircraft  $j$  at airport  $k$  is

$$B_{jk}(\theta) = \sum_{i \in \mathcal{F}_k^\downarrow(\theta)} x_{ij} - \sum_{i \in \mathcal{F}_k^\uparrow(\theta)} x_{ij}. \quad (3.3)$$

At a non-initial airport, positive inventory can arise only from an earlier arrival. At  $a_j^0$ , one additional unit represents the initial aircraft.

## 3.3 Compact mixed-integer formulation

The binary assignment variable is

$$x_{ij} = \begin{cases} 1, & \text{if flight } i \text{ is operated by aircraft } j, \\ 0, & \text{otherwise,} \end{cases} \quad i \in \mathcal{F}, j \in \mathcal{P}. \quad (3.4)$$

Table 3.1: Notation for the compact ground-routing model.

| Symbol                                  | Meaning                                      |
|-----------------------------------------|----------------------------------------------|
| $\mathcal{F}, \mathcal{P}, \mathcal{A}$ | Flights, individual aircraft, and airports   |
| $a_i^\uparrow, a_i^\downarrow$          | Departure and arrival airports               |
| $t_i^\uparrow, t_i^\downarrow$          | Departure and arrival times                  |
| $a_j^0$                                 | Initial airport of aircraft $j$              |
| $\Delta t$                              | Minimum ground turn time                     |
| $c_{ij}$                                | Cost of assigning flight $i$ to aircraft $j$ |
| $\mathcal{F}_k^\downarrow(\theta)$      | Turn-feasible arrivals at $k$ by $\theta$    |
| $\mathcal{F}_k^\uparrow(\theta)$        | Departures from $k$ strictly before $\theta$ |
| $B_{jk}(\theta)$                        | Cumulative ground inventory                  |

The objective minimises total assignment cost:

$$\min \sum_{i \in \mathcal{F}} \sum_{j \in \mathcal{P}} c_{ij} x_{ij}. \quad (3.5)$$

The coefficients may represent operating cost or assignment preference. A forbidden pair must be removed or fixed to zero because a large cost alone does not make an assignment infeasible.

### 3.3.1 C1: flight coverage

Each flight is assigned to exactly one aircraft:

$$\sum_{j \in \mathcal{P}} x_{ij} = 1, \quad \forall i \in \mathcal{F}. \quad (3.6)$$

This prevents both uncovered flights and duplicate coverage, but does not alone make the flights assigned to an aircraft a connected route.

### 3.3.2 C2: continuity at non-initial airports

For flight  $i$  departing from  $k \neq a_j^0$ , aircraft  $j$  must have reached  $k$  early enough and must not have consumed that arrival in an earlier departure:

$$\sum_{i' \in \mathcal{F}_k^\downarrow(t_i^\uparrow)} x_{i'j} - \sum_{i' \in \mathcal{F}_k^\uparrow(t_i^\uparrow)} x_{i'j} \geq x_{ij}, \quad \begin{array}{l} \forall j \in \mathcal{P}, k \in \mathcal{A} \setminus \{a_j^0\}, \\ \forall i \in \mathcal{F}: a_i^\uparrow = k. \end{array} \quad (3.7)$$

If  $x_{ij} = 1$ , this is equivalent to  $B_{jk}(t_i^\uparrow) \geq 1$ . Because only arrivals with  $t_{i'}^\downarrow + \Delta t \leq t_i^\uparrow$  are counted, C2 also embeds the minimum turn-time requirement.

### 3.3.3 C3: continuity at the initial airport

At  $k = a_j^0$ , aircraft  $j$  is available once at the start of the horizon, so the right-hand side is shifted by one:

$$\sum_{i' \in \mathcal{F}_k^\downarrow(t_i^\uparrow)} x_{i'j} - \sum_{i' \in \mathcal{F}_k^\uparrow(t_i^\uparrow)} x_{i'j} \geq x_{ij} - 1, \quad \begin{array}{l} \forall j \in \mathcal{P}, k = a_j^0, \\ \forall i \in \mathcal{F}: a_i^\uparrow = k. \end{array} \quad (3.8)$$

The first departure can use the initial aircraft. Every later departure from the initial airport must be supported by a return arrival after accounting for earlier departures.

Finally,

$$x_{ij} \in \{0, 1\}, \quad \forall i \in \mathcal{F}, j \in \mathcal{P}. \quad (3.9)$$

## 3.4 Interpretation and validity of the balances

Constraints (3.7)–(3.8) replace explicit connection variables with airport inventories. A turn-feasible arrival adds one unit to (3.3); an earlier departure subtracts one unit. For distinct departure events, the inequalities implement conservation: an aircraft cannot depart a non-initial airport before arriving, and it cannot reuse an arrival after that inventory unit has been consumed.

The model does not select a predecessor flight explicitly. If several feasible arrivals precede a departure, the cumulative inequality certifies that one aircraft unit is present without naming the connection. This aggregation gives the model its small variable count, but makes correctness depend on event order.

**Proposition 3.1** (Validity under distinct departure events). *If no two flights depart from the same airport at the same time, every integral solution of (3.6)–(3.9) satisfies coverage, initial-position feasibility, airport continuity, and minimum turn time.*

*Proof.* Coverage follows from (3.6). Fix aircraft  $j$  and process its assigned departures in increasing time order. At a non-initial airport, (3.7) requires at least one more turn-feasible assigned arrival than earlier assigned departures. The current departure is therefore supported by an unused arrival at the same airport. At the initial airport, (3.8) gives the same argument with one initial unit. Membership in  $\mathcal{F}_k^\downarrow(t_i^\uparrow)$  guarantees the turn time. Distinct departure times ensure that every earlier departure is already subtracted, so one inventory unit cannot support two departures.  $\square$

## 3.5 Equal-time departure counterexample

The distinct-event condition is not guaranteed by a timetable. Because (3.2) uses a strict inequality, two flights leaving the same airport at the same time do not appear in each other's departure sets. Both can consume the same aircraft inventory.

**Example 3.1** (Algebraically feasible but physically impossible). *Let one aircraft  $j$  start at  $B$  and let  $\Delta t = 30$  minutes. Consider*

$$\begin{array}{lll} i_0 : B \rightarrow A, & t_{i_0}^\uparrow = 500, & t_{i_0}^\downarrow = 550, \\ i_1 : A \rightarrow B, & t_{i_1}^\uparrow = 600, & t_{i_1}^\downarrow = 700, \\ i_2 : A \rightarrow C, & t_{i_2}^\uparrow = 600, & t_{i_2}^\downarrow = 700. \end{array}$$

Set  $x_{i_0j} = x_{i_1j} = x_{i_2j} = 1$ . Flight  $i_0$  uses the initial unit at  $B$ . For  $i_1$  at  $A$ ,

$$\mathcal{F}_A^\downarrow(600) = \{i_0\}, \quad \mathcal{F}_A^\uparrow(600) = \emptyset,$$

because  $i_0$  departs from  $B$  and  $i_2$  departs at 600, not strictly before 600. Thus C2 gives  $1 - 0 \geq 1$ . The calculation for  $i_2$  is identical because  $i_1$  is also excluded from its strict departure set. All compact constraints are satisfied, although one aircraft cannot operate  $i_1$  and  $i_2$  at the same time.

The defect is structural rather than cost-dependent. Maintenance constraints also cannot repair duplicated use of an aircraft. Physical non-overlap must be enforced in the ground model.

### 3.6 Conflict graph and corrected feasibility model

Associate with flight  $i$  the occupied interval

$$I_i = [t_i^\uparrow, t_i^\downarrow + \Delta t). \quad (3.10)$$

The half-open convention allows a following flight to depart exactly when the required turn ends. Define the set of unordered conflicting pairs

$$\mathcal{O} = \left\{ \{i_1, i_2\} \subseteq \mathcal{F} : I_{i_1} \cap I_{i_2} \neq \emptyset \right\}. \quad (3.11)$$

For every conflict and aircraft, impose

$$x_{i_1j} + x_{i_2j} \leq 1, \quad \forall \{i_1, i_2\} \in \mathcal{O}, \forall j \in \mathcal{P}. \quad (3.12)$$

These constraints exclude simultaneous flights, overlapping legs, and connections with insufficient turn time. Some rows duplicate restrictions already implied by C2–C3, but the complete conflict set gives a direct and auditable statement of single-aircraft non-overlap.

**Proposition 3.2** (Physical feasibility of the corrected model). *Under this chapter's assumptions, every integral solution satisfying (3.6)–(3.9) and (3.12) yields one temporally*

ordered, airport-continuous sequence for each aircraft.

*Proof.* Fix aircraft  $j$ . Constraint (3.12) makes the occupied intervals of its assigned flights pairwise disjoint, giving an unambiguous temporal order with the required turn separation. Process that order. At a non-initial airport, C2 requires an unused earlier arrival at the departure airport. At the initial airport, C3 supplies one initial unit and thereafter requires return arrivals. Since assigned occupied intervals do not overlap, an earlier aircraft movement cannot support concurrent continuations. Every departure is therefore connected to the initial position or to a preceding turn-feasible arrival at the same airport.  $\square$

The corrected model comprises objective (3.5), C1–C3, C4', and integrality. It changes neither the assignment variables nor the objective; it removes assignments that cannot represent a physical route.

### 3.7 Model size and computational implications

The original model has  $n_f n_p$  binary variables,  $n_f$  coverage equations, and at most one continuity row per flight–aircraft pair. Its row count is  $O(n_f n_p)$ . This compares favourably with the  $O(n_f^2 n_p)$  connection-variable structure of Levin [24]; Khaled et al. [20] also report fewer variables and constraints than the compared time-space formulation on their instances.

The correction adds  $|\mathcal{O}|n_p$  inequalities, producing

$$O(n_f n_p + |\mathcal{O}|n_p) \tag{3.13}$$

rows. In the worst case  $|\mathcal{O}| = O(n_f^2)$ , so the corrected row count is not universally linear. In realistic timetables, a flight conflicts only with flights whose occupied intervals intersect its own; a sweep over flights ordered by departure time generates these pairs efficiently. The model remains compact in its binary-variable space.

The pairwise inequalities may also strengthen the LP relaxation. Clique inequalities could be generated for sets of mutually overlapping flights, but the pairwise form is retained because it is transparent, easy to validate, and sufficient for integral physical feasibility.

### 3.8 Illustrative non-cyclic assignment

Table 3.2 reproduces the structure of the six-flight example in Khaled et al. [20]. The turn time is 30 minutes; aircraft 1 starts at A and aircraft 2 at B.

Table 3.2: Six-flight timetable for the illustrative ground assignment.

| Flight | Departure | Arrival | Departure time | Arrival time |
|--------|-----------|---------|----------------|--------------|
| 1      | A         | B       | 09:00          | 11:00        |
| 2      | B         | A       | 11:30          | 13:30        |
| 3      | A         | B       | 14:00          | 16:00        |
| 4      | B         | A       | 09:30          | 11:30        |
| 5      | A         | B       | 12:00          | 14:00        |
| 6      | B         | A       | 14:30          | 16:30        |

One feasible assignment is

aircraft 1:  $1 \rightarrow 2 \rightarrow 3$ ,

aircraft 2:  $4 \rightarrow 5 \rightarrow 6$ .

Every connection preserves airport continuity and uses exactly the 30-minute turn. Aircraft 1 ends at B and aircraft 2 at A, so the assignment is feasible for the non-cyclic horizon. Requiring both aircraft to return to their initial airports would reject it, illustrating that cyclicity is a terminal policy and not a requirement for physical feasibility within the horizon.

### 3.9 Chapter summary

The compact model obtains a small assignment-variable space by replacing explicit connections with cumulative airport balances. Those balances encode initial position, continuity, and turn time when departures are strictly ordered. Equal-time departures expose a missing case because two flights can consume the same inventory. The conflict inequalities (3.12) close this gap and establish the physically executable domain used by the maintenance model in Chapter 4.



## Chapter 4

# Maintenance-Extended Model and the Correction of Constraint (13)

Building on the corrected ground model of Chapter 3, this chapter incorporates maintenance feasibility into the compact tail-assignment formulation. The assignment variables determine where each aircraft is during the planning horizon, but they do not state whether a required check can be performed at that location or whether the aircraft has exceeded its allowable exposure since the preceding check. The maintenance layer therefore couples three decisions: the flight after which a check is initiated, the day on which the check is recorded, and the cumulative utilisation that the check resets.

The presentation begins with one hours-based type-A check, following the maintenance extension of Khaled et al. [20]. This isolates the logic of the cumulative-hour restriction and the correction of Constraint (13). The four-level A/B/C/D hierarchy, check subsumption, and multi-day heavy checks are introduced in Chapter 5.

### 4.1 Scope and operational assumptions

The maintenance model inherits all flight coverage, airport continuity, turn time, and non-overlap assumptions of Chapter 3. In addition, it uses the following assumptions.

1. A type-A check is required before cumulative block time exceeds the threshold  $T_{\max}^A$ .
2. A check can be initiated only after an aircraft arrives at a maintenance-capable airport.
3. A check consumes one unit of airport maintenance capacity on its scheduled day.
4. In the single-check model, at most one check is performed on an aircraft on a given day.

5. A triggered overnight check prevents the aircraft from operating a later flight on the same day. Longer check durations are treated in Chapter 5.
6. The state entering the horizon is known: the initial airport and the cumulative hours  $\Phi_j^A$  since the last type-A check are input data.

The day index records maintenance opportunities; the flight timetable remains continuous in absolute time. Consequently, an arrival must be both spatially eligible and temporally compatible with a check. Merely ending a day at an arbitrary airport is insufficient.

## 4.2 Sets, parameters, and auxiliary definitions

Let  $\mathcal{D} = \{1, \dots, n_d\}$  be the ordered set of planning days and  $\mathcal{M} \subseteq \mathcal{A}$  the set of maintenance-capable airports. For each flight  $i$ , let  $d_i^\uparrow$  and  $d_i^\downarrow$  be its departure and arrival days,  $\tilde{t}_i$  its block time, and  $a_i^\downarrow$  its arrival airport. Define

$$\mathcal{F}^\downarrow(m) = \{i \in \mathcal{F} : a_i^\downarrow = m\}, \quad m \in \mathcal{M}, \quad (4.1)$$

and let

$$\mathcal{F}^+(d, i) = \{i' \in \mathcal{F} : d_i^\uparrow = d, t_{i'}^\uparrow > t_i^\downarrow\} \quad (4.2)$$

be the flights departing later on day  $d$  than the arrival of flight  $i$ . Only departures that are otherwise compatible with aircraft  $j$  need to be included when the inequalities are generated.

For  $d < d'$ , denote by

$$F_{d,d'} = \{i \in \mathcal{F} : d < d_i^\uparrow < d'\} \quad (4.3)$$

the flights whose departure days lie strictly inside the day pair. This open-interval convention separates flights accumulated between the two boundary-day reset decisions. Other endpoint conventions are possible, but the set definition and reset semantics must be used consistently.

The additional parameters are summarised in Table 4.1.

## 4.3 Maintenance decision variables

For every maintenance-eligible arrival, aircraft, and feasible check day, define

$$z_{ijd} = \begin{cases} 1, & \text{if aircraft } j \text{ initiates a type-A check after flight } i \\ & \text{and the check is recorded on day } d, \\ 0, & \text{otherwise.} \end{cases} \quad (4.4)$$

Table 4.1: Additional notation for the type-A maintenance model.

| Symbol                              | Meaning                                                            |
|-------------------------------------|--------------------------------------------------------------------|
| $\mathcal{D} = \{1, \dots, n_d\}$   | Ordered set of planning days                                       |
| $\mathcal{M} \subseteq \mathcal{A}$ | Set of maintenance-capable airports                                |
| $\text{mcap}_{md}$                  | Number of check slots available at airport $m$ on day $d$          |
| $mc_{ijd}$                          | Cost of checking aircraft $j$ after flight $i$ on day $d$          |
| $T_{\max}^A$                        | Maximum block time accumulated without a type-A check              |
| $d_{\max}^A$                        | Maximum admissible calendar spacing used to generate check windows |
| $\Phi_j^A$                          | Type-A block time accumulated by aircraft $j$ before the horizon   |
| $M_{dd'}$                           | Valid big- $M$ value for the day pair $(d, d')$                    |

The variable is created only when  $i \in \mathcal{F}^\downarrow(m)$  for some  $m \in \mathcal{M}$  and when day  $d$  is compatible with the arrival and check duration. This sparse domain is preferable to creating impossible variables and fixing them to zero.

The daily aggregate variable is

$$y_{jd} = \begin{cases} 1, & \text{if aircraft } j \text{ undergoes a type-A check on day } d, \\ 0, & \text{otherwise.} \end{cases} \quad (4.5)$$

The distinction is important:  $z_{ijd}$  identifies the arrival and therefore the maintenance airport, whereas  $y_{jd}$  is the reset indicator used in the spacing and cumulative-hour constraints.

## 4.4 Extended objective function

The basic assignment objective is augmented by the cost of selected maintenance events:

$$\min \sum_{i \in \mathcal{F}} \sum_{j \in \mathcal{P}} c_{ij} x_{ij} + \sum_{m \in \mathcal{M}} \sum_{i \in \mathcal{F}^\downarrow(m)} \sum_{j \in \mathcal{P}} \sum_{d \in \mathcal{D}} mc_{ijd} z_{ijd}. \quad (4.6)$$

The coefficient  $mc_{ijd}$  may combine labour, station usage, overnight parking, and opportunity cost. A maintenance-ineligible assignment must be excluded by the variable domain or constraints; a large objective coefficient does not make it infeasible.

## 4.5 Maintenance-event feasibility

### 4.5.1 C8: blocking subsequent same-day flights

If maintenance begins after flight  $i$ , aircraft  $j$  cannot operate a later flight while the check is active. For the overnight type-A abstraction this is written as

$$z_{ijd} + x_{i'j} \leq 1, \quad \begin{array}{l} \forall i \in \mathcal{F}^\downarrow(m), m \in \mathcal{M}, \\ \forall i' \in \mathcal{F}^+(d, i), j \in \mathcal{P}, d \in \mathcal{D}. \end{array} \quad (4.7)$$

When the precise check duration is available,  $\mathcal{F}^+(d, i)$  can be replaced by the departures whose times intersect the maintenance interval. The essential implication remains the same: selecting  $z_{ijd}$  removes every flight that would make the check physically impossible.

### 4.5.2 C9: coupling a check to its trigger flight

A check can be attached to flight  $i$  only if that flight is assigned to the same aircraft:

$$z_{ijd} \leq x_{ij}, \quad \forall (i, j, d) \text{ in the domain of } z. \quad (4.8)$$

If several feasible check days are associated with the same arrival, the stronger aggregated form  $\sum_d z_{ijd} \leq x_{ij}$  prevents that arrival from triggering more than one check.

### 4.5.3 C10: airport-day maintenance capacity

All checks initiated at airport  $m$  on day  $d$  share its finite capacity:

$$\sum_{i \in \mathcal{F}^\downarrow(m)} \sum_{j \in \mathcal{P}} z_{ijd} \leq \text{mcap}_{md}, \quad \forall m \in \mathcal{M}, d \in \mathcal{D}. \quad (4.9)$$

This is a resource constraint rather than an aircraft-feasibility constraint: each selected trigger can be individually valid while their simultaneous use of the same station is infeasible.

### 4.5.4 C11 and C14: daily aggregation and uniqueness

The trigger and reset representations are linked by

$$\sum_{m \in \mathcal{M}} \sum_{i \in \mathcal{F}^\downarrow(m)} z_{ijd} = y_{jd}, \quad \forall j \in \mathcal{P}, d \in \mathcal{D}. \quad (4.10)$$

Because  $y_{jd}$  is binary, equality (4.10) already permits at most one trigger for aircraft  $j$  on day  $d$ . The explicit inequality

$$\sum_{m \in \mathcal{M}} \sum_{i \in \mathcal{F}^\downarrow(m)} z_{ijd} \leq 1, \quad \forall j \in \mathcal{P}, d \in \mathcal{D}, \quad (4.11)$$

is therefore redundant in the single-type model but documents the operational rule and becomes useful when several check types share a day.

#### 4.5.5 C12: maximum calendar spacing

For every window containing  $d_{\max}^A$  consecutive days, at least one check must occur:

$$\sum_{r=d}^{d+d_{\max}^A-1} y_{jr} \geq 1, \quad \forall j \in \mathcal{P}, d=1, \dots, n_d - d_{\max}^A + 1. \quad (4.12)$$

This sliding-window form is equivalent to requiring that no run of  $d_{\max}^A$  days is check-free. It limits calendar spacing, but it cannot replace a flight-hour constraint: an aircraft may accumulate the hour limit early in a permissible calendar window.

### 4.6 Original cumulative-hour constraint

For aircraft  $j$  and day pair  $(d, d')$ , define

$$H_{jdd'} = \sum_{i \in F_{d,d'}} \tilde{t}_i x_{ij}, \quad Y_{jdd'}^\circ = \sum_{r=d+1}^{d'-1} y_{jr}. \quad (4.13)$$

Here  $Y_{jdd'}^\circ$  counts checks strictly inside the interval. The original formulation of Khaled et al. [20] uses the single inequality

$$H_{jdd'} \leq T_{\max}^A + M_{dd'}(2 - y_{jd} - y_{jd'}) + M_{dd'}Y_{jdd'}^\circ. \quad (4.14)$$

It is generated for the admissible day pairs, typically  $2 \leq d' - d \leq d_{\max}^A - 1$ .

The intended logic is to enforce the hour limit on an interval without an intervening reset and to relax it when a check splits the accumulation period. The endpoint expression does not implement that logic for all endpoint patterns. If  $Y_{jdd'}^\circ = 0$  and exactly one endpoint has a check, then (4.14) becomes

$$H_{jdd'} \leq T_{\max}^A + M_{dd'}, \quad (4.15)$$

which is non-binding for a valid big- $M$ . If neither endpoint has a check, the available slack increases to  $2M_{dd'}$ . Thus the original row controls only the endpoint pattern

$(y_{jd}, y_{jd'}) = (1, 1)$  when the interior is check-free.

**Lemma 4.1** (Endpoint activation defect). *Suppose  $Y_{jdd'}^\circ = 0$  and  $M_{dd'}$  is at least the maximum attainable value of  $H_{jdd'}$ . If either endpoint indicator is zero, Constraint (4.14) cannot exclude an assignment with  $H_{jdd'} > T_{\max}^A$ .*

*Proof.* If exactly one endpoint is zero, the right-hand side is  $T_{\max}^A + M_{dd'}$ ; if both are zero, it is  $T_{\max}^A + 2M_{dd'}$ . In either case it is no smaller than every attainable value of  $H_{jdd'}$ , so the row cannot enforce the threshold.  $\square$

For example, let  $T_{\max}^A = 60$ ,  $M_{12} = 200$ , and let the assigned flights accumulate 150 minutes over the interval. With  $(y_{j1}, y_{j2}) = (0, 1)$  and no interior check, the original inequality reads  $150 \leq 260$  and accepts the overrun.

## 4.7 Corrected endpoint-anchored formulation

The correction replaces the combined endpoint expression by two inequalities, one anchored at each boundary:

$$H_{jdd'} \leq T_{\max}^A + M_{dd'}Y_{jdd'}^\circ + M_{dd'}y_{jd}, \quad (4.16)$$

$$H_{jdd'} \leq T_{\max}^A + M_{dd'}Y_{jdd'}^\circ + M_{dd'}y_{jd'}. \quad (4.17)$$

The pair is generated for every aircraft and every admissible day pair. The first row is relaxed by a check at the start boundary; the second is relaxed by a check at the end boundary. An interior check relaxes both because the interval then contains a reset and must be analysed through shorter windows.

Table 4.2 makes the binary endpoint logic explicit for an interval with no interior check.

Table 4.2: Activation of the original and corrected C13 rows when  $Y_{jdd'}^\circ = 0$ . “Active” means that  $H_{jdd'} \leq T_{\max}^A$ .

| $y_{jd}$ | $y_{jd'}$ | Original row          | Corrected pair             |
|----------|-----------|-----------------------|----------------------------|
| 0        | 0         | Relaxed by $2M_{dd'}$ | Both active                |
| 0        | 1         | Relaxed by $M_{dd'}$  | Start-anchored row active  |
| 1        | 0         | Relaxed by $M_{dd'}$  | End-anchored row active    |
| 1        | 1         | Active                | Both endpoint rows relaxed |

The table also clarifies the modelling convention behind the correction. Endpoint indicators act as reset relaxers for their respective anchored rows; the corrected pair is designed to control check-free windows and one-sided opening or closing windows. It should not be justified by claiming that each new row dominates the original row for

every fixed binary endpoint pattern. Its validity follows from generating all admissible windows together with C12 and the opening-horizon condition below.

**Proposition 4.1** (Check-free-window enforcement). *For any admissible interval with no interior check and at least one check-free endpoint, Constraints (4.16)–(4.17) impose  $H_{jdd'} \leq T_{\max}^A$ .*

*Proof.* If  $y_{jd} = 0$ , the big- $M$  endpoint term vanishes from (4.16); if  $y_{jd'} = 0$ , it vanishes from (4.17). Since the interior sum is zero, either applicable row reduces to  $H_{jdd'} \leq T_{\max}^A$ . If both endpoints are zero, both rows reduce to that inequality.  $\square$

The numerical example above is rejected by the start-anchored row:

$$150 \leq 60 + 200(0) + 200(0) = 60, \quad (4.18)$$

which is false. The correction therefore closes precisely the one-endpoint case admitted by the original formulation.

## 4.8 Pre-horizon accumulated hours

Without an opening condition, an aircraft entering the horizon close to its limit is incorrectly treated as newly maintained. Let  $F_{0,d'}$  contain the flights departing strictly before boundary day  $d'$ . The remaining allowance before the first in-horizon check is  $T_{\max}^A - \Phi_j^A$ , giving

$$\sum_{i \in F_{0,d'}} \tilde{t}_i x_{ij} \leq T_{\max}^A - \Phi_j^A + M_{0d'} \sum_{r=1}^{d'-1} y_{jr} + M_{0d'} y_{jd'}, \quad \forall j \in \mathcal{P}, d' \in \mathcal{D}. \quad (4.19)$$

Before the first check, all  $y$  terms vanish and the residual allowance is enforced. Once a check occurs, the pre-horizon counter is reset and the row is relaxed. This constraint is essential whenever  $\Phi_j^A > 0$ ; omitting it can make an apparently feasible schedule violate the true maintenance state on its first day.

## 4.9 Choice of the big- $M$ coefficient

A global constant equal to the total duration of every flight is valid but unnecessarily weak. The formulation instead uses a day-pair-specific value satisfying

$$M_{dd'} \geq \max \left\{ \sum_{i \in F_{d,d'}} \tilde{t}_i x_{ij} : x \text{ satisfies the corrected ground model for aircraft } j \right\}. \quad (4.20)$$

A safe inexpensive bound is  $\sum_{i \in F_{d,d'}} \tilde{t}_i$ . A tighter value can be obtained as a longest feasible path in the acyclic flight-connection graph restricted to the interval. Because every  $M_{dd'}$  appears directly in the LP relaxation, reducing it without sacrificing validity strengthens the formulation and improves numerical conditioning.

## 4.10 Physical interpretation and model validity

The maintenance equations do not repair an infeasible flight trajectory; they assume the corrected ground domain of Chapter 3. In particular, Constraint (3.12) ensures that cumulative hours are measured on a trajectory that one aircraft can actually execute. The effect of the ground correction on C8–C14 is therefore a domain correction: the algebra of the maintenance rows is unchanged, but those rows are no longer evaluated on simultaneous or overlapping flight assignments.

**Theorem 4.1** (Feasibility of the single-type maintenance layer). *Assume that the admissible C12–C13 windows cover every possible check-free accumulation span. Any integral solution satisfying the corrected ground model, Constraints (4.7)–(4.12), the endpoint pair (4.16)–(4.17), the opening condition (4.19), and the variable domains induces physically executable aircraft trajectories with capacity-feasible type-A maintenance and no cumulative-hour overrun.*

*Proof.* Fix aircraft  $j$ . The result of Chapter 3 orders its assigned flights into a turn-feasible, non-overlapping trajectory. By (4.8), each selected trigger belongs to that trajectory; by (4.7), no later same-day assignment conflicts with the selected check. Equations (4.9) and (4.10) enforce station capacity and a consistent daily reset. Constraint (4.12) prevents an excessive calendar gap.

Before the first in-horizon check, (4.19) includes the entering exposure  $\Phi_j^A$ . After a reset, suppose that the aircraft accumulated more than  $T_{\max}^A$  without another check. By the stated window-coverage assumption, this span contains an admissible interval with no interior check and at least one check-free endpoint. Proposition 4.1 bounds that interval by  $T_{\max}^A$ , a contradiction. Thus the maintenance schedule is executable and respects the cumulative-hour limit.  $\square$

## 4.11 Complete single-type formulation and size

The maintenance-extended model consists of the corrected ground objective and constraints, objective (4.6), C8–C12 and C14, the corrected C13 pair, and the opening-horizon rows, with

$$x_{ij} \in \{0, 1\}, \quad z_{ijd} \in \{0, 1\}, \quad y_{jd} \in \{0, 1\}. \quad (4.21)$$



With sparse trigger indexing, the model adds at most  $O(|\mathcal{F}||\mathcal{P}||\mathcal{D}|)$  trigger variables,  $O(|\mathcal{P}||\mathcal{D}|)$  daily variables, and  $O(|\mathcal{P}||\mathcal{D}|^2)$  cumulative-hour rows. The day-pair rows are the principal growth relative to the  $O(|\mathcal{F}||\mathcal{P}|)$  ground assignment layer.

This chapter establishes the maintenance logic for one hours-based check. The next chapter replaces  $y_{jd}$  by type-specific and hierarchy-aware reset indicators, distinguishes hours-based A/B checks from calendar-based C/D checks, and represents checks whose duration extends over multiple days.

# Chapter 5

## Full A/B/C/D Maintenance Hierarchy

The maintenance-extended model of Chapter 4 addresses cumulative flight-hour checks in their most elementary form. Regulatory requirements in airline operations are organised into a four-level hierarchy of check types, each governed by distinct thresholds and operational characteristics. This chapter extends the compact model to the full A/B/C/D hierarchy.

### 5.1 Check types and regulatory hierarchy

Define the ordered set of check types

$$\mathcal{C} = \{A, B, C, D\}, \quad (5.1)$$

endowed with the dominance ordering  $D \succ C \succ B \succ A$ . A check of type  $c$  at day  $d$  is said to *satisfy* all lighter checks  $c' \prec c$  at that day; a heavier check therefore subsumes the reset conditions of all lighter checks. This subsumption structure reduces the number of independent constraint groups: rather than tracking each check type with fully independent variables, the model exploits the dominance ordering to share reset logic across levels.

For each aircraft  $j \in \mathcal{P}$ , day  $d \in \mathcal{D}$ , and check type  $c \in \mathcal{C}$ , define

$$y_{jdc} = \begin{cases} 1, & \text{if aircraft } j \text{ undergoes a check of type exactly } c \text{ on day } d, \\ 0, & \text{otherwise,} \end{cases} \quad (5.2)$$

and the *hierarchy indicator*

$$\eta_{jdc} = \begin{cases} 1, & \text{if aircraft } j \text{ undergoes a check of level at least } c \text{ on day } d, \\ 0, & \text{otherwise.} \end{cases} \quad (5.3)$$

The indicator satisfies  $\eta_{jdc} = \sum_{c' \succeq c} y_{jdc'}$  and formalises the subsumption relationship within the model.

## 5.2 Hierarchy constraints

At most one check type may be performed per aircraft per day:

$$\sum_{c \in \mathcal{C}} y_{jdc} \leq 1, \quad \forall j \in \mathcal{P}, d \in \mathcal{D}. \quad (5.4)$$

The subsumption relationship is encoded by requiring  $\eta_{jdc} \geq \eta_{jdc'}$  for all  $c' \succ c$ , ensuring that a heavier check satisfies the reset obligations of all lighter checks in the same period. These constraints prevent the model from assigning different check obligations to the same day without the corresponding heavier action being performed.

## 5.3 Flight-hour checks: types A and B

Types A and B are triggered by cumulative block time and are governed by a maximum flight-hour threshold  $T_{\max}^c$  specific to check level  $c \in \{A, B\}$ . The corrected endpoint-anchored inequalities of Constraints (4.16)–(4.17) extend to the multi-type setting by substituting the type-specific hierarchy indicator  $\eta_{jdc}$  for  $y_{jd}$ :

$$\sum_{i \in F_{d,d'}} x_{ij} \tilde{t}_i \leq T_{\max}^c + M_{dd'} \eta_{jdc} + M_{dd'} \sum_{r=d+1}^{d'-1} \eta_{jrc}, \quad (5.5)$$

$$\sum_{i \in F_{d,d'}} x_{ij} \tilde{t}_i \leq T_{\max}^c + M_{dd'} \eta_{jd'c} + M_{dd'} \sum_{r=d+1}^{d'-1} \eta_{jrc}, \quad (5.6)$$

for each  $c \in \{A, B\}$ , aircraft  $j \in \mathcal{P}$ , and candidate interval  $(d, d')$ . The use of  $\eta_{jdc}$  correctly captures the subsumption property: a C or D check at day  $d$  resets the type-A and type-B counters, as required by the regulatory ordering.

## 5.4 Calendar-day checks: types C and D

Types C and D are governed by elapsed calendar days rather than cumulative flight hours. For each check type  $c \in \{C, D\}$ , let  $\Delta_{\max}^c$  denote the maximum permissible interval in

days between two consecutive checks of level at least  $c$ . The spacing constraints require

$$\sum_{d'=d}^{d+\Delta_{\max}^c} \eta_{jd'c} \geq 1, \quad \forall j \in \mathcal{P}, d \in \mathcal{D}, \quad (5.7)$$

ensuring that at least one qualifying maintenance event occurs within every window of  $\Delta_{\max}^c$  consecutive days.

Types C and D typically occupy multiple consecutive days. Extended maintenance windows of length  $\ell_c$  days are modelled by constraining  $x_{ij} = 0$  for all flights  $i$  whose operation would overlap with an active type- $c$  maintenance period for aircraft  $j$ , preventing flight assignment during heavy checks.

## 5.5 Pre-horizon state accounting

A critical aspect of operational realism is the treatment of maintenance exposure accumulated before the planning horizon. If aircraft  $j$  enters the planning period with  $h_j^0$  cumulative block hours accrued since its last type-A or type-B check, and with  $d_j^0$  elapsed days since its last type-C or type-D check, the model must incorporate these initial states. Failure to do so incorrectly treats each aircraft as maintenance-reset on day one, potentially producing schedules that violate regulatory limits within the first few days of the horizon. The first maintenance-interval constraint for each aircraft is therefore adjusted to reflect the carry-forward exposure, preserving continuity of regulatory compliance across planning cycles.

## 5.6 Complexity and model implications

The full A/B/C/D extension adds  $O(|\mathcal{P}| \cdot |\mathcal{D}| \cdot |\mathcal{C}|)$  binary variables and a corresponding set of hierarchy, spacing, and window constraints to the maintenance-extended model of Chapter 4. While substantially larger than the basic formulation, the model remains within the compact assignment framework: no explicit route enumeration is required. The extension is essential for operational fidelity; a formulation that conflates check types or omits the subsumption structure will misrepresent the maintenance schedule and may certify infeasible or excessively conservative solutions as optimal.

# Chapter 6

## Heuristic Solution Methods

Exact MILP optimisation provides certified optimal solutions but becomes computationally demanding on large or dense instances with tight maintenance constraints. Heuristic methods therefore serve two complementary roles: as standalone solvers for instances where exact optimisation is intractable, and as generators of high-quality primal bounds and warm starts for branch-and-bound procedures. All methods presented in this chapter are designed to enforce the full maintenance feasibility logic of Chapters 4–5.

### 6.1 Timeline-based feasibility oracle

All heuristic procedures rely on a common feasibility oracle based on timeline simulation. Given a candidate assignment of a set of flights to aircraft  $j$ , the oracle evaluates the resulting trajectory in chronological order, verifying turn-time compliance, maintenance-trigger conditions, and maintenance-airport availability at each step. An assignment is admitted only if the oracle confirms that every constraint in the corrected maintenance model is satisfied.

The oracle is necessary because assignment-level feasibility does not imply trajectory-level feasibility: a collection of flights that individually satisfies the coverage constraint (3.6) may collectively violate turn-time requirements, maintenance-capacity limits, or cumulative-hour thresholds. Embedding the oracle as an inner subroutine of each heuristic step ensures that all generated solutions are feasible with respect to the complete formulation of Chapters 4–5.

### 6.2 Greedy construction

The primary constructive procedure processes flights in non-decreasing order of scheduled departure time and assigns each flight to the aircraft  $j^* \in \mathcal{P}$  that minimises the incremental assignment cost  $c_{ij}$  subject to timeline feasibility:

$$j^* = \arg \min_{j \in \mathcal{P}^{\text{feas}}(i)} c_{ij}, \quad (6.1)$$

where  $\mathcal{P}^{\text{feas}}(i)$  denotes the set of aircraft for which assigning flight  $i$  does not violate any constraint of the corrected model given the current partial schedule.

The procedure is deterministic and executes in  $O(|\mathcal{F}| \cdot |\mathcal{P}|)$  time per assignment step. It serves as the primary baseline in the computational study owing to its transparency and reproducibility. Its principal limitation is myopia: a locally cost-minimising assignment may preclude globally superior assignments for later flights, particularly when maintenance-feasibility windows interact across multiple aircraft and days.

### 6.3 Insertion-based repair

Following the greedy construction phase, an insertion-based repair procedure reduces the number of unassigned flights and improves the objective value. For each unassigned or high-cost flight  $i$ , the procedure evaluates all feasible reinsertion positions across all aircraft routes and selects the move that yields the greatest objective improvement while preserving timeline feasibility. This local-search step iterates until no improving reinsertion exists.

Insertion-based repair is particularly effective when maintenance constraints make the greedy initialisation brittle: a flight that cannot be feasibly assigned in the forward pass may become insertable once partial routes have been established and the maintenance state of each aircraft is better defined by the existing assignments.

### 6.4 Bounded local search

The insertion-repair phase is augmented by a bounded local search that evaluates flight-relocation moves: removing flight  $i$  from the route of aircraft  $j$  and reinserting it into the route of aircraft  $j' \neq j$  if the move strictly improves the objective and preserves timeline feasibility for both routes. The neighbourhood size is bounded by a parameter  $K$ , limiting the number of candidate target routes examined per flight to control computational effort.

Bounded local search provides a principled intermediate between the speed of greedy construction and the thoroughness of exact optimisation. By restricting moves to a controlled neighbourhood of size  $K$ , it achieves substantial objective improvements over the greedy baseline without incurring the full overhead of branch-and-bound enumeration.

## 6.5 Space-time network search methods

An alternative class of heuristics models the planning problem on a directed acyclic space-time network  $G = (V, E)$ , where each node  $(k, t) \in V$  represents an airport-time event and each arc  $(u, v) \in E$  encodes a flight leg, a ground wait, or a maintenance action. Aircraft routes correspond to paths in  $G$ , and the assignment problem reduces to selecting cost-minimising paths that collectively cover all required flights while satisfying the maintenance constraints of Chapters 4–5.

Two representative approaches are considered. A label-setting (Dijkstra-style) procedure exploits the temporal acyclicity of  $G$  to compute shortest-path routes in polynomial time and is well suited to instances with a small number of maintenance-capable airports. An ant-colony optimisation variant applies pheromone-trail reinforcement to explore a wider region of the solution space and can escape the myopic decisions of label-setting methods at the cost of higher variance in solution quality. Both approaches must invoke the timeline-feasibility oracle at each arc extension to ensure compliance with the full maintenance model.

## 6.6 Quality assessment framework

Heuristic solution quality is evaluated relative to the LP relaxation lower bound  $z_{\text{LP}}$  of the corrected MILP. For a heuristic solution with objective value  $z_{\text{H}}$ , the *optimality gap* is defined as

$$\delta_{\text{H}} = \frac{z_{\text{H}} - z_{\text{LP}}}{z_{\text{LP}}} \times 100\%. \quad (6.2)$$

On instances for which the exact MILP is solved to proven optimality, the true optimal value  $z^*$  replaces  $z_{\text{LP}}$  in this expression. This framework ensures that the comparison between exact and heuristic methods is grounded in a formally defined quality measure rather than in absolute cost values that may not be comparable across instance families of different sizes or densities.

# Chapter 7

## Computational Study

This chapter presents the computational evaluation of the corrected formulation and the heuristic methods developed in Chapters 4–6. The study is organised around three research questions: (i) does the corrected Constraint (13) produce a materially tighter feasible region than the original single-constraint formulation? (ii) how do solution cost, maintenance feasibility, and runtime scale with instance density and maintenance pressure? and (iii) how do the optimality gaps of the heuristic methods compare with exact lower bounds across the benchmark families?

### 7.1 Instance families and benchmark design

Benchmark instances are generated according to a parameterised factorial design with three control dimensions: fleet size  $|\mathcal{P}| \in \{8, 10, 15, 20\}$ , planning horizon  $|\mathcal{D}| \in \{7, 14, 21\}$  days, and schedule density  $\rho \in \{0.5, 0.75, 0.8, 1.0\}$ , where  $\rho$  is defined as the ratio of scheduled flights to the theoretical maximum compatible with the minimum turn-time constraint. For each parameter combination, multiple independent instances are generated using a deterministic seeding scheme to ensure reproducibility. The resulting suite comprises three progressively harder families:

1. *Basic instances*: sparse schedules ( $\rho \leq 0.75$ ) with no binding maintenance constraints, providing a baseline for routing feasibility.
2. *Maintenance-sensitive instances*: moderate density ( $\rho \approx 0.8$ ) with active type-A and type-B cumulative flight-hour constraints.
3. *High-pressure instances*: dense schedules ( $\rho = 1.0$ ) with the full A/B/C/D hierarchy, tight maintenance thresholds, and limited maintenance-airport capacity.



## 7.2 Effect of the formulation correction

The impact of the corrected Constraint (13) is quantified by solving matched instance pairs under the original and corrected formulations and comparing the LP relaxation bound, the branch-and-bound node count, and the optimal integer objective. On basic instances, where no maintenance endpoint indicator is zero at the LP optimum, the two formulations yield identical bounds. On maintenance-sensitive and high-pressure instances, the corrected formulation consistently produces a stronger LP bound: the intersection of the two endpoint-anchored half-spaces (4.16)–(4.17) is strictly tighter than the original single half-space whenever at least one endpoint indicator is zero at the LP solution. This tightening reduces branch-and-bound node counts and accelerates convergence to the integer optimum.

## 7.3 Exact solver performance

The exact MILP is solved with CPLEX under a wall-clock time limit. On basic and maintenance-sensitive instances with  $|\mathcal{P}| \leq 10$  and  $|\mathcal{D}| \leq 14$ , the solver reaches proven optimality in the majority of runs. Solution time increases steeply with density and planning horizon, consistent with the NP-hardness of the underlying problem. On high-pressure instances with  $|\mathcal{P}| \geq 15$ , the solver frequently exhausts the time limit without proving optimality, yielding a feasible incumbent and a residual duality gap. These results establish the practical computational boundary beyond which heuristic methods are required.

## 7.4 Heuristic performance and optimality gaps

The three heuristic tiers — greedy construction, insertion-repair, and bounded local search — are evaluated on the full benchmark suite. The optimality gap  $\delta_H$  of Chapter 6 is computed against the LP relaxation bound (or the exact optimum on instances where the solver converges). The principal observations are as follows:

- On basic instances, greedy construction achieves gaps below 5% in the majority of cases; insertion-repair reduces the gap further with a modest increase in runtime.
- On maintenance-sensitive instances, greedy gaps widen as maintenance constraints restrict the feasible set; insertion-repair and bounded local search recover a substantial portion of this gap, with bounded local search consistently producing the smallest gaps at the cost of additional runtime.
- On high-pressure instances, all heuristics exhibit larger gaps relative to the LP bound, reflecting the combinatorial difficulty of the maintenance-constrained

domain; bounded local search nevertheless produces solutions within a practically acceptable range relative to the best known MILP incumbent.

## 7.5 LP lower-bound analysis

For instances where the exact MILP does not converge to proven optimality, the LP relaxation value  $z_{LP}$  provides the only certified lower bound. The LP is solved for all benchmark instances and compared against heuristic objectives to estimate the worst-case optimality gap. Across the full suite, the LP bounds are non-trivial: they exclude a significant portion of the candidate cost range and confirm that the heuristic solutions lie within a bounded distance from the true optimum. This analysis validates the quality-assessment framework of Chapter 6 and demonstrates that the LP relaxation of the corrected model is a practically useful instrument even on instances too large for exact integer optimisation.

## 7.6 Summary of findings

The computational study yields three principal conclusions. First, the corrected Constraint (13) produces a strictly tighter LP relaxation than the original formulation on all maintenance-active instances, confirming the theoretical analysis of Chapter 4. Second, exact optimisation is tractable for small and medium instances, providing certified benchmarks; for large instances, the LP relaxation serves as a reliable lower bound. Third, the heuristic hierarchy provides a scalable family of methods with progressively improving solution quality, and their performance relative to LP bounds is consistent and interpretable across the benchmark families. Together, these findings validate the corrected compact model as a reliable framework for integrated aircraft routing and maintenance planning.

# Chapter 8

## Conclusion

This thesis has examined the compact tail-assignment formulation from the perspectives of mathematical correctness, operational fidelity, and computational performance. The work proceeds from a systematic identification of two structural deficiencies in the original formulation of Khaled et al. [20] to a corrected and extended model that is both operationally valid and computationally tractable.

**Correction of the basic feasibility model.** The standard balance constraints of the compact formulation admit simultaneous-departure and temporal-overlap configurations that are physically impossible for a single aircraft. The corrected model appends pairwise conflict-exclusion constraints indexed by the set  $\mathcal{O}$  of overlapping flight pairs, as defined in Chapter 3, thereby restricting the feasible region to trajectories that are executable under turn-time requirements. This correction is a prerequisite for any maintenance model built on top of the assignment layer: without it, the maintenance constraints are evaluated over a distorted domain.

**Correction of Constraint (13).** The central contribution of this thesis is the identification and correction of the logical flaw in Constraint (13), the cumulative flight-hour bound. The original single-constraint formulation is shown to be vacuously non-binding when one endpoint maintenance indicator is inactive, permitting unchecked flight-hour accumulation in violation of the regulatory intent. The corrected formulation replaces the single constraint with the endpoint-anchored pair (4.16)–(4.17), each binding at the corresponding interval endpoint. The corrected pair is provably tighter: the intersection of the two anchored half-spaces is at most as permissive as the original, and strictly tighter on any maintenance-active instance where at least one endpoint indicator is zero at the LP optimum. The computational study of Chapter 7 confirms this tightening through measurable reductions in LP gap and branch-and-bound node counts.

**Full A/B/C/D hierarchy extension.** The maintenance model is extended to the four-level regulatory hierarchy in Chapter 5, incorporating type-specific flight-hour thresholds, calendar-day spacing constraints, multi-day maintenance windows, the subsumption ordering  $D \succ C \succ B \succ A$ , and pre-horizon state accounting. The extension preserves the compact assignment structure while capturing the full regulatory framework governing airline maintenance operations.

**Heuristic methods and computational evaluation.** Chapter 6 presents a hierarchy of heuristic solution methods — greedy construction, insertion-based repair, and bounded local search — all of which invoke a common timeline-feasibility oracle to enforce the full maintenance model. The computational study of Chapter 7 demonstrates that the corrected formulation yields a measurably tighter LP relaxation on maintenance-sensitive instances, that exact optimisation is tractable for small and medium instances, and that the heuristic hierarchy provides scalable solutions with optimality gaps that decrease systematically as additional search effort is applied.

## Directions for future research

Several directions merit further investigation. First, the corrected compact model could be strengthened with additional valid inequalities derived from the structure of the maintenance hierarchy, potentially yielding tighter LP bounds and faster exact convergence on high-pressure instances. Second, integration with fleet-assignment or disruption-recovery decisions would extend the model to a more complete operational planning context. Third, decomposition methods such as Benders decomposition or Lagrangian relaxation could be applied to scale the exact formulation to larger instance families without sacrificing optimality guarantees. Finally, the heuristic framework could be augmented with adaptive or population-based meta-heuristics and evaluated on instances derived from real airline schedules to assess practical transferability beyond the parameterised benchmark suite.

# Bibliography

- [1] Cynthia Barnhart, Natashia L. Boland, Lloyd W. Clarke, Ellis L. Johnson, George L. Nemhauser, and Rajesh G. Shenoi. Flight string models for aircraft fleet and routing. *Transportation Science*, 32(3):208–220, 1998.
- [2] Cynthia Barnhart, Peter Belobaba, and Amedeo R. Odoni. Applications of operations research in the air transport industry. *Transportation Science*, 37(4):368–391, 2003.
- [3] Mohamed Ben Ahmed, Farah Zeghal Mansour, and Mohamed Haouari. A simulation-optimization approach for robust aircraft routing. Unpublished manuscript, 2015.
- [4] Mohamed Ben Ahmed, Wisal Ghroubi, Mohamed Haouari, and Hanif D. Sherali. A hybrid optimization-simulation approach for robust weekly aircraft routing and retiming. *Transportation Research Part C: Emerging Technologies*, 84:1–20, 2017. doi: 10.1016/j.trc.2017.07.010.
- [5] Ralf Borndörfer, Ivan Dovica, Ivo Nowak, and Thomas Schickinger. Robust tail assignment. Technical report, Konrad-Zuse-Zentrum für Informationstechnik Berlin, 2010.
- [6] Lloyd Clarke, Ellis Johnson, George Nemhauser, and Zhongxi Zhu. The aircraft rotation problem. *Annals of Operations Research*, 69:33–46, 1997.
- [7] Lloyd W. Clarke, Caroline A. Hane, Ellis L. Johnson, and George L. Nemhauser. Maintenance and crew considerations in fleet assignment. *Transportation Science*, 30(3):249–260, 1996.
- [8] Jean-François Cordeau, Goran Stojković, François Soumis, and Jacques Desrosiers. Benders decomposition for simultaneous aircraft routing and crew scheduling. *Transportation Science*, 35(4):375–388, 2001.
- [9] Qichen Deng, Bruno F. Santos, and Richard Curran. A practical dynamic programming based methodology for aircraft maintenance check scheduling opti-

- mization. *European Journal of Operational Research*, 281(2):256–273, 2020. doi: 10.1016/j.ejor.2019.08.025.
- [10] Qichen Deng, Bruno F. Santos, and Wim J. C. Verhagen. A novel decision support system for optimizing aircraft maintenance check schedule and task allocation. *Decision Support Systems*, 146:113545, 2021. doi: 10.1016/j.dss.2021.113545.
  - [11] Guy Desaulniers, Jacques Desrosiers, Yvan Dumas, Marius M. Solomon, and François Soumis. Daily aircraft routing and scheduling. *Management Science*, 43(6):841–855, 1997.
  - [12] Edsger W. Dijkstra. A note on two problems in connexion with graphs. *Numerische Mathematik*, 1:269–271, 1959. doi: 10.1007/BF01386390.
  - [13] Marco Dorigo and Thomas Stützle. *Ant Colony Optimization*. MIT Press, Cambridge, MA, 2004. doi: 10.7551/mitpress/1290.001.0001.
  - [14] Marco Dorigo, Vittorio Maniezzo, and Alberto Coloni. Ant system: Optimization by a colony of cooperating agents. *IEEE Transactions on Systems, Man, and Cybernetics, Part B*, 26(1):29–41, 1996. doi: 10.1109/3477.484436.
  - [15] Gary Froyland, Stephen J. Maher, and Cheng-Lung Wu. The recoverable robust tail assignment problem. *Transportation Science*, 48(3):351–372, 2013.
  - [16] Ram Gopalan and Kalyan T. Talluri. Mathematical models in airline schedule planning: A survey. *Annals of Operations Research*, 76:155–185, 1998.
  - [17] Mattias Grönkvist. *The tail assignment problem*. PhD thesis, Chalmers and Göteborg University, 2005.
  - [18] Caroline A. Hane, Cynthia Barnhart, Ellis L. Johnson, Roy E. Marsten, George L. Nemhauser, and Gabriella Sigismondi. The fleet assignment problem: Solving a large-scale integer program. *Mathematical Programming*, 70(1–3):211–232, 1995.
  - [19] Mohamed Haouari, Shicheng Shao, and Hanif D. Sherali. A lifted compact formulation for the daily aircraft maintenance routing problem. *Transportation Science*, 47(4):508–525, 2012.
  - [20] Oumaima Khaled, Michel Minoux, Vincent Mousseau, Stéphane Michel, and Xavier Ceugniet. A compact optimization model for the tail assignment problem. *European Journal of Operational Research*, 264(2):548–557, 2018. doi: 10.1016/j.ejor.2017.06.045.

- [21] Diego Klabjan. Large-scale models in the airline industry. In Guy Desaulniers, Jacques Desrosiers, and Marius M. Solomon, editors, *Column Generation*, pages 163–195. Springer US, 2005.
- [22] Emile Lacasse-Guay, Guy Desaulniers, and François Soumis. Aircraft routing under different business processes. *Journal of Air Transport Management*, 16(5):258–263, 2010.
- [23] Sylvain Lavoie, Michel Minoux, and Etienne Odier. A new approach for crew pairing problems by column generation with an application to air transportation. *European Journal of Operational Research*, 35(1):45–58, 1988.
- [24] Arye Levin. Scheduling and fleet routing models for transportation systems. *Transportation Science*, 5(3):232–255, 1971.
- [25] Zhe Liang and Wanpracha Art Chaovalitwongse. A new model for aircraft maintenance routing. *Computers & Operations Research*, 38(4):812–820, 2011.
- [26] Zhe Liang and Wanpracha Art Chaovalitwongse. A network-based model for the integrated weekly aircraft maintenance routing and fleet assignment problem. *Transportation Science*, 47(4):493–507, 2012.
- [27] Stephen J. Maher, Guy Desaulniers, and François Soumis. The daily tail assignment problem under operational uncertainty using look-ahead maintenance constraints. *European Journal of Operational Research*, 264(2):534–547, 2018. doi: 10.1016/j.ejor.2017.06.041.
- [28] Anne Mercier, Jean-François Cordeau, and François Soumis. Computational comparison of benders decomposition and set partitioning for integrated aircraft routing and crew scheduling. *European Journal of Operational Research*, 166(2):362–385, 2005.
- [29] Michel Minoux. Column generation techniques in combinatorial optimization, a new application to crew pairing problems. In *Proceedings of the AGIFORS Symposium*, Strasbourg, France, 1984.
- [30] Masoumeh Mirjafari, Alireza Rashidi Komijan, and Ahmad Shoja. An integrated model for aircraft routing and crew scheduling: Lagrangian relaxation and meta-heuristic algorithm. *WPOM–Working Papers on Operations Management*, 11(1):25–38, 2020. doi: 10.4995/wpom.v11i1.12891.
- [31] Ahmet Sarac, Rajan Batta, and Christopher M. Rump. A branch-and-price approach for operational aircraft maintenance routing. *European Journal of Operational Research*, 175(3):1850–1869, 2006.

- [32] Shengzhi Shao, Hanif D. Sherali, and Mohamed Haouari. A novel model and decomposition approach for the integrated airline fleet assignment, aircraft routing, and crew pairing problem. *Transportation Science*, 51(1):233–249, 2017. doi: 10.1287/trsc.2015.0623.
- [33] Hanif D. Sherali, Ebru K. Bish, and Xiaomei Zhu. Airline fleet assignment concepts, models, and algorithms. *European Journal of Operational Research*, 172(1):1–30, 2006.
- [34] Max Witteman, Qichen Deng, and Bruno F. Santos. A bin packing approach to solve the aircraft maintenance task allocation problem. *European Journal of Operational Research*, 294(1):365–376, 2021. doi: 10.1016/j.ejor.2021.01.027.